

Version: Free, Pro, & Team ▼







GitHub Actions / Reference / Workflows and actions / Workflow syntax

Workflow syntax for GitHub Actions

A workflow is a configurable automated process made up of one or more jobs. You must create a YAML file to define your workflow configuration.

In this article About YAML syntax for workflows name run-name on on.<event_name>.types on.<pull_request|pull_request_target>.<branches|branches-ignore> on.push.
branches|tags|branches-ignore|tags-ignore> on.<push|pull_request|pull_request_target>.<paths|paths-ignore> on.schedule on.workflow_call on.workflow_call.inputs on.workflow_call.inputs.<input_id>.type on.workflow_call.outputs on.workflow_call.secrets on.workflow_call.secrets.<secret_id> on.workflow_call.secrets.<secret_id>.required on.workflow_run.
branches|branches-ignore> on.workflow_dispatch on.workflow_dispatch.inputs on.workflow_dispatch.inputs.<input_id>.required on.workflow_dispatch.inputs.<input_id>.type permissions

How permissions are calculated for a workflow job

env

defaults

defaults.run

defaults.run.shell

defaults.run.working-directory

concurrency

jobs

jobs.<job_id>

jobs.<job_id>.name

jobs.<job_id>.permissions

jobs.<job_id>.needs

jobs.<job_id>.if

jobs.<job_id>.runs-on

jobs.<job_id>.environment

jobs.<job_id>.concurrency

jobs.<job_id>.outputs

jobs.<job_id>.env

jobs.<job_id>.defaults

jobs.<job_id>.defaults.run

jobs.<job_id>.defaults.run.shell

jobs.<job_id>.defaults.run.working-directory

jobs.<job_id>.steps

jobs.<job_id>.steps[*].id

jobs.<job_id>.steps[*].if

jobs.<job_id>.steps[*].name

jobs.<job_id>.steps[*].uses

jobs.<job_id>.steps[*].run

jobs.<job_id>.steps[*].working-directory

jobs.<job_id>.steps[*].shell

jobs.<job_id>.steps[*].with

jobs.<job_id>.steps[*].with.args

jobs.<job_id>.steps[*].with.entrypoint

jobs.<job_id>.steps[*].env

jobs.<job_id>.steps[*].continue-on-error

jobs.<job_id>.steps[*].timeout-minutes

```
jobs.<job_id>.timeout-minutes
```

jobs.<job_id>.strategy

jobs.<job_id>.strategy.matrix

jobs.<job_id>.strategy.matrix.include

jobs.<job_id>.strategy.matrix.exclude

jobs.<job_id>.strategy.fail-fast

jobs.<job_id>.strategy.max-parallel

jobs.<job_id>.continue-on-error

jobs.<job_id>.container

jobs.<job_id>.container.image

jobs.<job_id>.container.credentials

jobs.<job_id>.container.env

jobs.<job_id>.container.ports

jobs.<job_id>.container.volumes

jobs.<job_id>.container.options

jobs.<job_id>.services

jobs.<job_id>.services.<service_id>.image

jobs.<job_id>.services.<service_id>.credentials

jobs.<job_id>.services.<service_id>.env

jobs.<job_id>.services.<service_id>.ports

jobs.<job_id>.services.<service_id>.volumes

jobs.<job_id>.services.<service_id>.options

jobs.<job_id>.uses

jobs.<job_id>.with

jobs.<job_id>.with.<input_id>

jobs.<job_id>.secrets

jobs.<job_id>.secrets.inherit

jobs.<job_id>.secrets.<secret_id>

Filter pattern cheat sheet

About YAML syntax for workflows ∂

Workflow files use YAML syntax, and must have either a .yml or .yaml file extension. If you're new to YAML and want to learn more, see Learn YAML in Y minutes.

You must store workflow files in the _github/workflows directory of your repository.

name ∂

The name of the workflow. GitHub displays the names of your workflows under your repository's "Actions" tab. If you omit name, GitHub displays the workflow file path relative to the root of the repository.

run-name ∂

The name for workflow runs generated from the workflow. GitHub displays the workflow run name in the list of workflow runs on your repository's "Actions" tab. If run-name is omitted or is only whitespace, then the run name is set to event-specific information for the workflow run. For example, for a workflow triggered by a push or pull_request event, it is set as the commit message or the title of the pull request.

This value can include expressions and can reference the github and inputs contexts.

Example of run-name &

```
run-name: Deploy to ${{ inputs.deploy_target }} by @${{ github.actor }}
```

on Ø

To automatically trigger a workflow, use on to define which events can cause the workflow to run. For a list of available events, see Events that trigger workflows.

You can define single or multiple events that can trigger a workflow, or set a time schedule. You can also restrict the execution of a workflow to only occur for specific files, tags, or branch changes. These options are described in the following sections.

Using a single event &

For example, a workflow with the following on value will run when a push is made to any branch in the workflow's repository:

```
on: push
```

Using multiple events *∂*

You can specify a single event or multiple events. For example, a workflow with the following on value will run when a push is made to any branch in the repository or when someone forks the repository:

```
on: [push, fork]
```

If you specify multiple events, only one of those events needs to occur to trigger your workflow. If multiple triggering events for your workflow occur at the same time, multiple workflow runs will be triggered.

Using activity types *∂*

Some events have activity types that give you more control over when your workflow should run. Use on.<event_name>.types to define the type of event activity that will trigger a workflow run.

For example, the <code>issue_comment</code> event has the <code>created</code>, <code>edited</code>, and <code>deleted</code> activity types. If your workflow triggers on the <code>label</code> event, it will run whenever a label is created, edited, or deleted. If you specify the <code>created</code> activity type for the <code>label</code> event, your workflow will run when a label is created but not when a label is edited or deleted.

```
on:
label:
types:
- created
```

If you specify multiple activity types, only one of those event activity types needs to occur to trigger your workflow. If multiple triggering event activity types for your workflow occur at the same time, multiple workflow runs will be triggered. For example, the following workflow triggers when an issue is opened or labeled. If an issue with two labels is opened, three workflow runs will start: one for the issue opened event and two for the two issue labeled events.

```
on:
issues:
types:
- opened
- labeled
```

For more information about each event and their activity types, see Events that trigger workflows.

Using filters *⊘*

Some events have filters that give you more control over when your workflow should run.

For example, the push event has a branches filter that causes your workflow to run only when a push to a branch that matches the branches filter occurs, instead of when any push occurs.

```
on:
   push:
   branches:
   - main
   - 'releases/**'
```

Using activity types and filters with multiple events *⊘*

If you specify activity types or filters for an event and your workflow triggers on multiple events, you must configure each event separately. You must append a colon (:) to all events, including events without configuration.

For example, a workflow with the following on value will run when:

- · A label is created
- A push is made to the main branch in the repository
- A push is made to a GitHub Pages-enabled branch

```
on:
    label:
    types:
        - created
    push:
        branches:
```

- main
page_build:

on.<event_name>.types ∂

Use on.<event_name>.types to define the type of activity that will trigger a workflow run. Most GitHub events are triggered by more than one type of activity. For example, the label is triggered when a label is created, edited, or deleted. The types keyword enables you to narrow down activity that causes the workflow to run. When only one activity type triggers a webhook event, the types keyword is unnecessary.

You can use an array of event types . For more information about each event and their activity types, see Events that trigger workflows.

on:
label:
types: [created, edited]

on.<pull_request|pull_request_target>. <branches|branches-ignore> ∂

When using the pull_request and pull_request_target events, you can configure a workflow to run only for pull requests that target specific branches.

Use the branches filter when you want to include branch name patterns or when you want to both include and exclude branch names patterns. Use the branches—ignore filter when you only want to exclude branch name patterns. You cannot use both the branches and branches—ignore filters for the same event in a workflow.

If you define both branches / branches-ignore and paths / paths-ignore , the workflow will only run when both filters are satisfied.

The branches and branches—ignore keywords accept glob patterns that use characters like *, **, +, ?, ! and others to match more than one branch name. If a name contains any of these characters and you want a literal match, you need to escape each of these special characters with \. For more information about glob patterns, see the Workflow syntax for GitHub Actions.

Example: Including branches *∂*

The patterns defined in branches are evaluated against the Git ref's name. For example, the following workflow would run whenever there is a pull_request event for a pull request targeting:

- A branch named main (refs/heads/main)
- A branch named mona/octocat (refs/heads/mona/octocat)
- A branch whose name starts with releases/, like releases/10 (refs/heads/releases/10)

```
on:
   pull_request:
     # Sequence of patterns matched against refs/heads
     branches:
     - main
     - 'mona/octocat'
     - 'releases/**'
```

If a workflow is skipped due to branch filtering, <u>path filtering</u>, or a <u>commit message</u>, then checks associated with that workflow will remain in a "Pending" state. A pull request that requires those checks to be successful will be blocked from merging.

Example: Excluding branches \mathscr{O}

When a pattern matches the branches—ignore pattern, the workflow will not run. The patterns defined in branches—ignore are evaluated against the Git ref's name. For example, the following workflow would run whenever there is a pull_request event unless the pull request is targeting:

- A branch named mona/octocat (refs/heads/mona/octocat)
- A branch whose name matches releases/**-alpha, like releases/beta/3-alpha (refs/heads/releases/beta/3-alpha)

```
on:
   pull_request:
    # Sequence of patterns matched against refs/heads
   branches-ignore:
    - 'mona/octocat'
    - 'releases/**-alpha'
```

Example: Including and excluding branches &

You cannot use branches and branches—ignore to filter the same event in a single workflow. If you want to both include and exclude branch patterns for a single event, use the branches filter along with the ! character to indicate which branches should be excluded.

If you define a branch with the ! character, you must also define at least one branch without the ! character. If you only want to exclude branches, use branches—ignore instead.

The order that you define patterns matters.

- A matching negative pattern (prefixed with !) after a positive match will exclude the Git ref.
- A matching positive pattern after a negative match will include the Git ref again.

The following workflow will run on <code>pull_request</code> events for pull requests that target releases/10 or releases/beta/mona, but not for pull requests that target releases/10-alpha or releases/beta/3-alpha because the negative pattern <code>!releases/**-alpha</code> follows the positive pattern.

```
on:
   pull_request:
    branches:
    - 'releases/**'
    - '!releases/**-alpha'
```

on.push.
 branches|tags|branches-ignore|tags-ignore> ∂

When using the push event, you can configure a workflow to run on specific branches or tags.

Use the branches filter when you want to include branch name patterns or when you want to both include and exclude branch names patterns. Use the branches—ignore filter when you only want to exclude branch name patterns. You cannot use both the branches and branches—ignore filters for the same event in a workflow.

Use the tags filter when you want to include tag name patterns or when you want to both include and exclude tag names patterns. Use the tags-ignore filter when you only want to exclude tag name patterns. You cannot use both the tags and tags-ignore filters for the same event in a workflow.

If you define only tags / tags-ignore or only branches / branches-ignore, the workflow won't run for events affecting the undefined Git ref. If you define neither tags / tags-ignore or branches / branches-ignore, the workflow will run for events affecting either branches or tags. If

you define both branches / branches-ignore and paths / paths-ignore , the workflow will only run when both filters are satisfied.

The branches, branches-ignore, tags, and tags-ignore keywords accept glob patterns that use characters like *, **, +, ?, ! and others to match more than one branch or tag name. If a name contains any of these characters and you want a literal match, you need to escape each of these special characters with \. For more information about glob patterns, see the Workflow syntax for GitHub Actions.

Example: Including branches and tags *∂*

The patterns defined in branches and tags are evaluated against the Git ref's name. For example, the following workflow would run whenever there is a push event to:

- A branch named main (refs/heads/main)
- A branch named mona/octocat (refs/heads/mona/octocat)
- A branch whose name starts with releases/, like releases/10 (refs/heads/releases/10)
- A tag named v2 (refs/tags/v2)
- A tag whose name starts with v1., like v1.9.1 (refs/tags/v1.9.1)

Example: Excluding branches and tags *⊘*

When a pattern matches the branches—ignore or tags—ignore pattern, the workflow will not run. The patterns defined in branches and tags are evaluated against the Git ref's name. For example, the following workflow would run whenever there is a push event, unless the push event is to:

A branch named mona/octocat (refs/heads/mona/octocat)

- A branch whose name matches releases/**-alpha, like releases/beta/3-alpha (refs/heads/releases/beta/3-alpha)
- A tag named v2 (refs/tags/v2)
- A tag whose name starts with v1., like v1.9 (refs/tags/v1.9)

Example: Including and excluding branches and tags *⊘*

You can't use branches and branches—ignore to filter the same event in a single workflow. Similarly, you can't use tags and tags—ignore to filter the same event in a single workflow. If you want to both include and exclude branch or tag patterns for a single event, use the branches or tags filter along with the ! character to indicate which branches or tags should be excluded.

If you define a branch with the ! character, you must also define at least one branch without the ! character. If you only want to exclude branches, use branches—ignore instead. Similarly, if you define a tag with the ! character, you must also define at least one tag without the ! character. If you only want to exclude tags, use tags—ignore instead.

The order that you define patterns matters.

- A matching negative pattern (prefixed with !) after a positive match will exclude the Git ref.
- A matching positive pattern after a negative match will include the Git ref again.

The following workflow will run on pushes to releases/10 or releases/beta/mona, but not on releases/10-alpha or releases/beta/3-alpha because the negative pattern !releases/**-alpha follows the positive pattern.

```
on:
  push:
  branches:
```

```
- 'releases/**'
- '!releases/**-alpha'
```

on.<push|pull_request|pull_request_target>. <paths|paths=ignore> ₽

When using the push and pull_request events, you can configure a workflow to run based on what file paths are changed. Path filters are not evaluated for pushes of tags.

Use the paths filter when you want to include file path patterns or when you want to both include and exclude file path patterns. Use the paths-ignore filter when you only want to exclude file path patterns. You cannot use both the paths and paths-ignore filters for the same event in a workflow. If you want to both include and exclude path patterns for a single event, use the paths filter prefixed with the ! character to indicate which paths should be excluded.

Note

The order that you define paths patterns matters:

- A matching negative pattern (prefixed with !) after a positive match will exclude the path.
- A matching positive pattern after a negative match will include the path again.

If you define both branches / branches-ignore and paths / paths-ignore , the workflow will only run when both filters are satisfied.

The paths and paths-ignore keywords accept glob patterns that use the * and ** wildcard characters to match more than one path name. For more information, see the Workflow syntax for GitHub Actions.

Example: Including paths ∂

If at least one path matches a pattern in the paths filter, the workflow runs. For example, the following workflow would run anytime you push a JavaScript file (.js).

```
on:
  push:
  paths:
  - '**.js'
```

If a workflow is skipped due to path filtering, <u>branch filtering</u>, or a <u>commit message</u>, then checks associated with that workflow will remain in a "Pending" state. A pull request that requires those checks to be successful will be blocked from merging.

Example: Excluding paths *∂*

When all the path names match patterns in paths-ignore, the workflow will not run. If any path names do not match patterns in paths-ignore, even if some path names match the patterns, the workflow will run.

A workflow with the following path filter will only run on push events that include at least one file outside the docs directory at the root of the repository.

```
on:
   push:
   paths-ignore:
   - 'docs/**'
```

Example: Including and excluding paths *∂*

You cannot use paths and paths-ignore to filter the same event in a single workflow. If you want to both include and exclude path patterns for a single event, use the paths filter prefixed with the ! character to indicate which paths should be excluded.

If you define a path with the ! character, you must also define at least one path without the ! character. If you only want to exclude paths, use paths-ignore instead.

The order that you define paths patterns matters:

- A matching negative pattern (prefixed with !) after a positive match will exclude the path.
- A matching positive pattern after a negative match will include the path again.

This example runs anytime the push event includes a file in the sub-project directory or its subdirectories, unless the file is in the sub-project/docs directory. For example, a push that changed sub-project/index.js or sub-project/src/index.js will trigger a workflow run, but a push changing only sub-project/docs/readme.md will not.

```
on:
push:
paths:
```

- 'sub-project/**'
- '!sub-project/docs/**'

Git diff comparisons *∂*

(i) Note

If you push more than 1,000 commits, or if GitHub does not generate the diff due to a timeout, the workflow will always run.

The filter determines if a workflow should run by evaluating the changed files and running them against the paths-ignore or paths list. If there are no files changed, the workflow will not run.

GitHub generates the list of changed files using two-dot diffs for pushes and three-dot diffs for pull requests:

- **Pull requests:** Three-dot diffs are a comparison between the most recent version of the topic branch and the commit where the topic branch was last synced with the base branch.
- Pushes to existing branches: A two-dot diff compares the head and base SHAs directly with each other.
- Pushes to new branches: A two-dot diff against the parent of the ancestor of the deepest commit pushed.

Note

Diffs are limited to 300 files. If there are files changed that aren't matched in the first 300 files returned by the filter, the workflow will not run. You may need to create more specific filters so that the workflow will run automatically.

For more information, see About comparing branches in pull requests.

on.schedule @

You can use on schedule to define a time schedule for your workflows. You can schedule a workflow to run at specific UTC times using POSIX cron syntax. Scheduled workflows run on the latest commit on the default or base branch. The shortest interval you can run scheduled workflows is once every 5 minutes.

This example triggers the workflow every day at 5:30 and 17:30 UTC:

```
on:
    schedule:
    # * is a special character in YAML so you have to quote this string
    - cron: '30 5,17 * * *'
```

A single workflow can be triggered by multiple schedule events. You can access the schedule event that triggered the workflow through the github event schedule context. This example triggers the workflow to run at 5:30 UTC every Monday-Thursday, but skips the Not on Monday or Wednesday step on Monday and Wednesday.

For more information about cron syntax, see Events that trigger workflows.

on.workflow_call ∂

Use on workflow_call to define the inputs and outputs for a reusable workflow. You can also map the secrets that are available to the called workflow. For more information on reusable workflows, see Reuse workflows.

on.workflow_call.inputs ∂

When using the workflow_call keyword, you can optionally specify inputs that are passed to the called workflow from the caller workflow. For more information about the workflow_call keyword, see Events that trigger workflows.

In addition to the standard input parameters that are available, on.workflow_call.inputs requires a type parameter. For more information, see on.workflow_call.inputs input_id>.type.

If a default parameter is not set, the default value of the input is false for a boolean, 0 for a number, and "" for a string.

Within the called workflow, you can use the inputs context to refer to an input. For more information, see Contexts reference.

If a caller workflow passes an input that is not specified in the called workflow, this results in an error.

Example of on.workflow_call.inputs &

```
on:
    workflow_call:
        inputs:
        username:
            description: 'A username passed from the caller workflow'
            default: 'john-doe'
            required: false
            type: string

jobs:
    print-username:
    runs-on: ubuntu-latest

    steps:
        - name: Print the input name to STDOUT
            run: echo The username is ${{ inputs.username }}
```

For more information, see Reuse workflows.

on.workflow_call.inputs.<input_id>.type ∂

Required if input is defined for the on.workflow_call keyword. The value of this parameter is a string specifying the data type of the input. This must be one of: boolean, number, or string.

on.workflow_call.outputs ∂

A map of outputs for a called workflow. Called workflow outputs are available to all downstream jobs in the caller workflow. Each output has an identifier, an optional description, and a value. The value must be set to the value of an output from a job within the called workflow.

In the example below, two outputs are defined for this reusable workflow: workflow_output1 and workflow_output2. These are mapped to outputs called job_output1 and job_output2, both from a job called my_job.

Example of on.workflow_call.outputs &

```
on:
    workflow_call:
    # Map the workflow outputs to job outputs
    outputs:
        workflow_output1:
        description: "The first job output"
        value: ${{ jobs.my_job.outputs.job_output1 }}
    workflow_output2:
        description: "The second job output"
        value: ${{ jobs.my_job.outputs.job_output2 }}
```

For information on how to reference a job output, see <a href="jobs.<job_id>.outputs">jobs.<job_id>.outputs. For more information, see Reuse workflows.

on.workflow_call.secrets ℰ

A map of the secrets that can be used in the called workflow.

Within the called workflow, you can use the secrets context to refer to a secret.

Note

If you are passing the secret to a nested reusable workflow, then you must use jobs. job_id>.secrets again to pass the secret. For more information, see Reuse workflows.

If a caller workflow passes a secret that is not specified in the called workflow, this results in an error.

Example of on.workflow_call.secrets @

```
on:
 workflow call:
    secrets:
      access-token:
        description: 'A token passed from the caller workflow'
        required: false
jobs:
  pass-secret-to-action:
    runs-on: ubuntu-latest
    steps:
    # passing the secret to an action
      - name: Pass the received secret to an action
        uses: ./.github/actions/my-action
          token: ${{ secrets.access-token }}
  # passing the secret to a nested reusable workflow
  pass-secret-to-workflow:
    uses: ./.github/workflows/my-workflow
    secrets:
       token: ${{ secrets.access-token }}
```

on.workflow_call.secrets.<secret_id> @

A string identifier to associate with the secret.

on.workflow_call.secrets.<secret_id>.required @

A boolean specifying whether the secret must be supplied.

on.workflow_run.
branches|branches-ignore> ∂

When using the workflow_run event, you can specify what branches the triggering workflow must run on in order to trigger your workflow.

The branches and branches—ignore filters accept glob patterns that use characters like *,

**, +, ?, ! and others to match more than one branch name. If a name contains any of these
characters and you want a literal match, you need to escape each of these special characters with

\text{\characters} For more information about glob patterns, see the Workflow syntax for GitHub Actions.

For example, a workflow with the following trigger will only run when the workflow named Build runs on a branch whose name starts with releases/:

```
on:
  workflow_run:
    workflows: ["Build"]
  types: [requested]
  branches:
    - 'releases/**'
```

A workflow with the following trigger will only run when the workflow named Build runs on a branch that is not named canary:

```
on:
  workflow_run:
    workflows: ["Build"]
    types: [requested]
    branches-ignore:
    - "canary"
```

You cannot use both the branches and branches—ignore filters for the same event in a workflow. If you want to both include and exclude branch patterns for a single event, use the branches filter along with the ! character to indicate which branches should be excluded.

The order that you define patterns matters.

- A matching negative pattern (prefixed with !) after a positive match will exclude the branch.
- A matching positive pattern after a negative match will include the branch again.

For example, a workflow with the following trigger will run when the workflow named Build runs on a branch that is named releases/10 or releases/beta/mona but will not releases/10—alpha, releases/beta/3—alpha, or main.

```
on:
    workflow_run:
    workflows: ["Build"]
    types: [requested]
    branches:
        - 'releases/**'
        - '!releases/**-alpha'
```

on.workflow_dispatch ∂

When using the workflow_dispatch event, you can optionally specify inputs that are passed to the workflow.

This trigger only receives events when the workflow file is on the default branch.

on.workflow_dispatch.inputs ∂

The triggered workflow receives the inputs in the inputs context. For more information, see Contexts.

(i) Note

- The workflow will also receive the inputs in the github.event.inputs context. The information in the inputs context and github.event.inputs context is identical except that the inputs context preserves Boolean values as Booleans instead of converting them to strings. The choice type resolves to a string and is a single selectable option.
- The maximum number of top-level properties for inputs is 10.
- The maximum payload for inputs is 65,535 characters.

Example of on.workflow_dispatch.inputs &

```
on:
 workflow_dispatch:
    inputs:
      logLevel:
        description: 'Log level'
        required: true
        default: 'warning'
        type: choice
        options:
          - info
          warning
          debug
      print_tags:
        description: 'True to print to STDOUT'
        required: true
        type: boolean
```

```
description: 'Test scenario tags'
    required: true
    type: string
    environment:
        description: 'Environment to run tests against'
        type: environment
        required: true

jobs:
    print-tag:
    runs-on: ubuntu-latest
    if: ${{ inputs.print_tags }}
    steps:
        - name: Print the input tag to STDOUT
        run: echo The tags are ${{ inputs.tags }}
```

on.workflow_dispatch.inputs.<input_id>.required



A boolean specifying whether the input must be supplied.

on.workflow_dispatch.inputs.<input_id>.type ∂

The value of this parameter is a string specifying the data type of the input. This must be one of: boolean, choice, number, environment or string.

permissions ∂

You can use permissions to modify the default permissions granted to the GITHUB_TOKEN, adding or removing access as required, so that you only allow the minimum required access. For more information, see Use GITHUB_TOKEN for authentication in workflows.

You can use permissions either as a top-level key, to apply to all jobs in the workflow, or within specific jobs. When you add the permissions key within a specific job, all actions and run commands within that job that use the GITHUB_TOKEN gain the access rights you specify. For more information, see jobs.<job_id>.permissions.

Owners of an organization can restrict write access for the GITHUB_TOKEN at the repository level. For more information, see Disabling or limiting GitHub Actions for your organization.

When a workflow is triggered by the pull_request_target event, the GITHUB_TOKEN is granted read/write repository permission, even when it is triggered from a public fork. For more information, see Events that trigger workflows.

For each of the available permissions, shown in the table below, you can assign one of the access levels: read (if applicable), write, or none. write includes read. If you specify the access for any of these permissions, all of those that are not specified are set to none.

Available permissions and details of what each allows an action to do:

Permission	Allows an action using GITHUB_TOKEN to
actions	Work with GitHub Actions. For example, actions: write permits an action to cancel a workflow run.
	For more information, see Permissions required for GitHub Apps.
attestations	Work with artifact attestations. For example, attestations: write permits an action to
	generate an artifact attestation for a build. For more information, see Using artifact attestations to establish provenance for builds
checks	Work with check runs and check suites. For example, checks: write permits an action to create a check run. For more information, see
	Permissions required for GitHub Apps.
contents	Work with the contents of the repository. For example, contents: read permits an action to list the commits, and contents: write allows the
	action to create a release. For more information, see Permissions required for GitHub Apps.
deployments	Work with deployments. For example, deployments: write permits an action to create a
	new deployment. For more information, see Permissions required for GitHub Apps.
discussions	Work with GitHub Discussions. For example, discussions: write permits an action to close or
	delete a discussion. For more information, see <u>Using</u> the GraphQL API for Discussions.

Permission	Allows an action using GITHUB_TOKEN to
id-token	Fetch an OpenID Connect (OIDC) token. This requires id-token: write. For more information, see OpenID Connect
issues	Work with issues. For example, issues: write permits an action to add a comment to an issue. For more information, see Permissions required for GitHub Apps.
models	Generate AI inference responses with GitHub Models. For example, models: read permits an action to use the GitHub Models inference API. See Prototyping with AI models.
packages	Work with GitHub Packages. For example, packages: write permits an action to upload and publish packages on GitHub Packages. For more information, see About permissions for GitHub Packages.
pages	Work with GitHub Pages. For example, pages: write permits an action to request a GitHub Pages build. For more information, see Permissions required for GitHub Apps.
pull-requests	Work with pull requests. For example, pull-requests: write permits an action to add a label to a pull request. For more information, see Permissions required for GitHub Apps.
security-events	Work with GitHub code scanning alerts. For example, security-events: read permits an action to list the code scanning alerts for the repository, and security-events: write allows an action to update the status of a code scanning alert. For more information, see Repository permissions for 'Code scanning alerts'.
	Dependabot and secret scanning alerts cannot be read with this permission and require a GitHub App or a personal access token. For more information, see Repository permissions for 'Dependabot alerts'

Permission	Allows an action using GITHUB_TOKEN to
	and Repository permissions for 'Secret scanning
	alerts' in "Permissions required for GitHub Apps."
statuses	Work with commit statuses. For example,
	statuses: read permits an action to list the commit
	statuses for a given reference. For more information,
	see Permissions required for GitHub Apps.

Defining access for the **GITHUB_TOKEN** scopes *∂*

You can define the access that the GITHUB_TOKEN will permit by specifying read, write, or none as the value of the available permissions within the permissions key.

```
permissions:
    actions: read|write|none
    attestations: read|write|none
    checks: read|write|none
    contents: read|write|none
    deployments: read|write|none
    id-token: write|none
    issues: read|write|none
    models: read|none
    discussions: read|write|none
    packages: read|write|none
    pages: read|write|none
    pull-requests: read|write|none
    security-events: read|write|none
    statuses: read|write|none
```

If you specify the access for any of these permissions, all of those that are not specified are set to none.

You can use the following syntax to define one of read-all or write-all access for all of the available permissions:

```
permissions: read-all
```

https://docs.github.com/en/actions/reference/workflows-and-actions/workflow-syntax

permissions: write-all

You can use the following syntax to disable permissions for all of the available permissions:

```
permissions: {}
```

Changing the permissions in a forked repository \mathscr{O}

You can use the permissions key to add and remove read permissions for forked repositories, but typically you can't grant write access. The exception to this behavior is where an admin user has selected the **Send write tokens to workflows from pull requests** option in the GitHub Actions settings. For more information, see Managing GitHub Actions settings for a repository.

How permissions are calculated for a workflow job ∂

The permissions for the GITHUB_TOKEN are initially set to the default setting for the enterprise, organization, or repository. If the default is set to the restricted permissions at any of these levels then this will apply to the relevant repositories. For example, if you choose the restricted default at the organization level then all repositories in that organization will use the restricted permissions as the default. The permissions are then adjusted based on any configuration within the workflow file, first at the workflow level and then at the job level. Finally, if the workflow was triggered by a pull request from a forked repository, and the **Send write tokens to workflows from pull requests** setting is not selected, the permissions are adjusted to change any write permissions to read only.

Setting the **GITHUB_TOKEN** permissions for all jobs in a workflow *⊘*

You can specify permissions at the top level of a workflow, so that the setting applies to all jobs in the workflow.

Example: Setting the GITHUB_TOKEN permissions for an entire workflow &

This example shows permissions being set for the GITHUB_TOKEN that will apply to all jobs in the workflow. All permissions are granted read access.

```
name: "My workflow"

on: [ push ]

permissions: read-all
```

jobs:

Using the **permissions** key for forked repositories $\mathscr E$

You can use the permissions key to add and remove read permissions for forked repositories, but typically you can't grant write access. The exception to this behavior is where an admin user has selected the **Send write tokens to workflows from pull requests** option in the GitHub Actions settings. For more information, see Managing GitHub Actions settings for a repository.

Permissions for workflow runs triggered by Dependabot *⊘*

Workflow runs triggered by Dependabot pull requests run as if they are from a forked repository, and therefore use a read-only GITHUB_TOKEN. These workflow runs cannot access any secrets. For information about strategies to keep these workflows secure, see Secure use reference.

env 2

A map of variables that are available to the steps of all jobs in the workflow. You can also set variables that are only available to the steps of a single job or to a single step. For more information, see jobs.<job_id>.env and jobs.<job_id>.steps[*].env.

Variables in the env map cannot be defined in terms of other variables in the map.

When more than one environment variable is defined with the same name, GitHub uses the most specific variable. For example, an environment variable defined in a step will override job and workflow environment variables with the same name, while the step executes. An environment variable defined for a job will override a workflow variable with the same name, while the job executes.

Example of env ?

env:

SERVER: production

defaults ∂

Use defaults to create a map of default settings that will apply to all jobs in the workflow. You can also set default settings that are only available to a job. For more information, see jobs.
job_id>.defaults.

When more than one default setting is defined with the same name, GitHub uses the most specific default setting. For example, a default setting defined in a job will override a default setting that has the same name defined in a workflow.

defaults.run ∂

You can use defaults.run to provide default shell and working-directory options for all run steps in a workflow. You can also set default settings for run that are only available to a job. For more information, see jobs.<job_id>.defaults.run. You cannot use contexts or expressions in this keyword.

When more than one default setting is defined with the same name, GitHub uses the most specific default setting. For example, a default setting defined in a job will override a default setting that has the same name defined in a workflow.

Example: Set the default shell and working directory &

defaults:
 run:
 shell: bash
 working-directory: ./scripts

defaults.run.shell ∂

Use shell to define the shell for a step. This keyword can reference several contexts. For more information, see Contexts.

Supported platform	shell parameter	Description	Command run internally
Linux / macOS	unspecified	The default shell on non-Windows platforms. Note that this runs a different command to when bash is specified	bash -e {0}

Supported platform	shell parameter	Description	Command run internally
		explicitly. If bash is not found in the path, this is treated as sh.	
All	bash	The default shell on non-Windows platforms with a fallback to sh. When specifying a bash shell on Windows, the bash shell included with Git for Windows is used.	<pre>bashnoprofile norc -eo pipefail {0}</pre>
All	pwsh	The PowerShell Core. GitHub appends the extension .ps1 to your script name.	<pre>pwsh -command ". '{0}'"</pre>
All	python	Executes the python command.	python {0}
Linux / macOS	sh	The fallback behavior for non-Windows platforms if no shell is provided and bash is not found in the path.	sh -e {0}
Windows	cmd	GitHub appends the extension .cmd to your script name and substitutes for {0}.	%ComSpec% /D /E:ON /V:OFF /S /C "CALL " {0}"".
Windows	pwsh	This is the default shell used on Windows. The PowerShell Core. GitHub appends the extension .ps1 to your script name. If your self-hosted Windows runner does not have PowerShell Core installed, then	<pre>pwsh -command ". '{0}'".</pre>

Supported platform	shell parameter	Description	Command run internally
		PowerShell Desktop is used instead.	
Windows	powershell	The PowerShell Desktop. GitHub appends the extension .ps1 to your script name.	<pre>powershell -command ". '{0}'".</pre>

When more than one default setting is defined with the same name, GitHub uses the most specific default setting. For example, a default setting defined in a job will override a default setting that has the same name defined in a workflow.

defaults.run.working-directory ∂

Use working-directory to define the working directory for the shell for a step. This keyword can reference several contexts. For more information, see Contexts.



Ensure the working-directory you assign exists on the runner before you run your shell in it. When more than one default setting is defined with the same name, GitHub uses the most specific default setting. For example, a default setting defined in a job will override a default setting that has the same name defined in a workflow.

concurrency ∂

Use concurrency to ensure that only a single job or workflow using the same concurrency group will run at a time. A concurrency group can be any string or expression. The expression can only use github, inputs and vars contexts. For more information about expressions, see Evaluate expressions in workflows and actions.

You can also specify concurrency at the job level. For more information, see jobs. <job_id>.concurrency.

This means that there can be at most one running and one pending job in a concurrency group at any time. When a concurrent job or workflow is queued, if another job or workflow using the same concurrency group in the repository is in progress, the queued job or workflow will be pending.

Any existing pending job or workflow in the same concurrency group, if it exists, will be canceled and the new queued job or workflow will take its place.

To also cancel any currently running job or workflow in the same concurrency group, specify cancel—in—progress: true. To conditionally cancel currently running jobs or workflows in the same concurrency group, you can specify cancel—in—progress as an expression with any of the allowed expression contexts.

Note

- The concurrency group name is case insensitive. For example, prod and Prod will be treated as the same concurrency group.
- Ordering is not guaranteed for jobs or workflow runs using concurrency groups. Jobs or workflow runs in the same concurrency group are handled in an arbitrary order.

Example: Using concurrency and the default behavior @

The default behavior of GitHub Actions is to allow multiple jobs or workflow runs to run concurrently. The concurrency keyword allows you to control the concurrency of workflow runs.

For example, you can use the concurrency keyword immediately after where trigger conditions are defined to limit the concurrency of entire workflow runs for a specific branch:

```
on:
   push:
    branches:
    - main

concurrency:
   group: ${{ github.workflow }}-${{ github.ref }}
   cancel-in-progress: true
```

You can also limit the concurrency of jobs within a workflow by using the concurrency keyword at the job level:

```
on:
   push:
    branches:
        - main

jobs:
   job-1:
   runs-on: ubuntu-latest
```

```
concurrency:
group: example-group
cancel-in-progress: true
```

Example: Concurrency groups $\mathscr O$

Concurrency groups provide a way to manage and limit the execution of workflow runs or jobs that share the same concurrency key.

The concurrency key is used to group workflows or jobs together into a concurrency group. When you define a concurrency key, GitHub Actions ensures that only one workflow or job with that key runs at any given time. If a new workflow run or job starts with the same concurrency key, GitHub Actions will cancel any workflow or job already running with that key. The concurrency key can be a hard-coded string, or it can be a dynamic expression that includes context variables.

It is possible to define concurrency conditions in your workflow so that the workflow or job is part of a concurrency group.

This means that when a workflow run or job starts, GitHub will cancel any workflow runs or jobs that are already in progress in the same concurrency group. This is useful in scenarios where you want to prevent parallel runs for a certain set of a workflows or jobs, such as the ones used for deployments to a staging environment, in order to prevent actions that could cause conflicts or consume more resources than necessary.

In this example, job-1 is part of a concurrency group named staging_environment. This means that if a new run of job-1 is triggered, any runs of the same job in the staging_environment concurrency group that are already in progress will be cancelled.

```
jobs:
   job-1:
    runs-on: ubuntu-latest
    concurrency:
        group: staging_environment
        cancel-in-progress: true
```

Alternatively, using a dynamic expression such as concurrency: ci-\${{ github.ref }} in your workflow means that the workflow or job would be part of a concurrency group named ci-followed by the reference of the branch or tag that triggered the workflow. In this example, if a new commit is pushed to the main branch while a previous run is still in progress, the previous run will be cancelled and the new one will start:

```
on:
    push:
    branches:
    - main

concurrency:
    group: ci-${{ github.ref }}
    cancel-in-progress: true
```

Example: Using concurrency to cancel any in-progress job or run *∂*

To use concurrency to cancel any in-progress job or run in GitHub Actions, you can use the concurrency key with the cancel-in-progress option set to true:

```
concurrency:
  group: ${{ github.ref }}
  cancel-in-progress: true
```

Note that in this example, without defining a particular concurrency group, GitHub Actions will cancel *any* in-progress run of the job or workflow.

Example: Using a fallback value *∂*

If you build the group name with a property that is only defined for specific events, you can use a fallback value. For example, <code>github.head_ref</code> is only defined on <code>pull_request</code> events. If your workflow responds to other events in addition to <code>pull_request</code> events, you will need to provide a fallback to avoid a syntax error. The following concurrency group cancels in-progress jobs or runs on <code>pull_request</code> events only; if <code>github.head_ref</code> is undefined, the concurrency group will fallback to the run ID, which is guaranteed to be both unique and defined for the run.

```
concurrency:
  group: ${{ github.head_ref || github.run_id }}
  cancel-in-progress: true
```

Example: Only cancel in-progress jobs or runs for the current workflow *∂*

If you have multiple workflows in the same repository, concurrency group names must be unique across workflows to avoid canceling in-progress jobs or runs from other workflows. Otherwise, any previously in-progress or pending job will be canceled, regardless of the workflow.

To only cancel in-progress runs of the same workflow, you can use the github.workflow property to build the concurrency group:

```
concurrency:
  group: ${{ github.workflow }}-${{ github.ref }}
  cancel-in-progress: true
```

Example: Only cancel in-progress jobs on specific branches *∂*

If you would like to cancel in-progress jobs on certain branches but not on others, you can use conditional expressions with cancel-in-progress. For example, you can do this if you would like to cancel in-progress jobs on development branches but not on release branches.

To only cancel in-progress runs of the same workflow when not running on a release branch, you can set cancel-in-progress to an expression similar to the following:

```
concurrency:
  group: ${{ github.workflow }}-${{ github.ref }}
  cancel-in-progress: ${{ !contains(github.ref, 'release/')}}
```

In this example, multiple pushes to a release/1.2.3 branch would not cancel in-progress runs. Pushes to another branch, such as main, would cancel in-progress runs.

jobs ∂

A workflow run is made up of one or more <code>jobs</code>, which run in parallel by default. To run jobs sequentially, you can define dependencies on other jobs using the <code>jobs.<job_id>.needs</code> keyword.

Each job runs in a runner environment specified by runs-on.

You can run an unlimited number of jobs as long as you are within the workflow usage limits. For more information, see <u>Billing and usage</u> for GitHub-hosted runners and <u>Actions limits</u> for self-hosted runner usage limits.

If you need to find the unique identifier of a job running in a workflow run, you can use the GitHub API. For more information, see REST API endpoints for GitHub Actions.

jobs.<job_id> ₽

Use <code>jobs.<job_id></code> to give your job a unique identifier. The key <code>job_id</code> is a string and its value is a map of the job's configuration data. You must replace <code><job_id></code> with a string that is unique to the <code>jobs</code> object. The <code><job_id></code> must start with a letter or <code>_</code> and contain only alphanumeric characters, <code>-</code>, or <code>_</code>.

Example: Creating jobs &

In this example, two jobs have been created, and their job_id values are my_first_job and my_second_job.

```
jobs:
   my_first_job:
    name: My first job
   my_second_job:
    name: My second job
```

jobs.<job_id>.name ♂

Use jobs.<job_id>.name to set a name for the job, which is displayed in the GitHub UI.

jobs.<job_id>.permissions ∂

For a specific job, you can use <code>jobs.<job_id>.permissions</code> to modify the default permissions granted to the <code>GITHUB_TOKEN</code>, adding or removing access as required, so that you only allow the minimum required access. For more information, see <code>Use GITHUB_TOKEN</code> for authentication in workflows.

By specifying the permission within a job definition, you can configure a different set of permissions for the GITHUB_TOKEN for each job, if required. Alternatively, you can specify the permissions for all jobs in the workflow. For information on defining permissions at the workflow level, see permissions.

For each of the available permissions, shown in the table below, you can assign one of the access levels: read (if applicable), write, or none. write includes read. If you specify the access for any of these permissions, all of those that are not specified are set to none.

Available permissions and details of what each allows an action to do:

Permission	Allows an action using GITHUB_TOKEN to
actions	Work with GitHub Actions. For example, actions: write permits an action to cancel a workflow run. For more information, see Permissions required for GitHub Apps.
attestations	Work with artifact attestations. For example, attestations: write permits an action to generate an artifact attestation for a build. For more information, see Using artifact attestations to establish provenance for builds
checks	Work with check runs and check suites. For example, checks: write permits an action to create a check run. For more information, see Permissions required for GitHub Apps.
contents	Work with the contents of the repository. For example, contents: read permits an action to list the commits, and contents: write allows the action to create a release. For more information, see Permissions required for GitHub Apps.
deployments	Work with deployments. For example, deployments: write permits an action to create a new deployment. For more information, see Permissions required for GitHub Apps.
discussions	Work with GitHub Discussions. For example, discussions: write permits an action to close or delete a discussion. For more information, see <u>Using</u> the GraphQL API for Discussions.
id-token	Fetch an OpenID Connect (OIDC) token. This requires id-token: write. For more information, see OpenID Connect
issues	Work with issues. For example, issues: write permits an action to add a comment to an issue. For more information, see Permissions required for GitHub Apps.

Permission	Allows an action using GITHUB_TOKEN to	
models	Generate AI inference responses with GitHub Models. For example, models: read permits an action to use the GitHub Models inference API. See Prototyping with AI models.	
packages	Work with GitHub Packages. For example, packages: write permits an action to upload and publish packages on GitHub Packages. For more information, see About permissions for GitHub Packages.	
pages	Work with GitHub Pages. For example, pages: write permits an action to request a GitHub Pages build. For more information, see Permissions required for GitHub Apps.	
pull-requests	Work with pull requests. For example, pull-requests: write permits an action to add a label to a pull request. For more information, see Permissions required for GitHub Apps.	
security-events	Work with GitHub code scanning alerts. For example, security-events: read permits an action to list the code scanning alerts for the repository, and security-events: write allows an action to update the status of a code scanning alert. For more information, see Repository permissions for 'Code scanning alerts'.	
	Dependabot and secret scanning alerts cannot be read with this permission and require a GitHub App or a personal access token. For more information, see Repository permissions for 'Dependabot alerts' and Repository permissions for 'Secret scanning alerts' in "Permissions required for GitHub Apps."	
statuses	Work with commit statuses. For example, statuses: read permits an action to list the commit statuses for a given reference. For more information, see Permissions required for GitHub Apps.	

Defining access for the **GITHUB_TOKEN** scopes *∂*

You can define the access that the GITHUB_TOKEN will permit by specifying read, write, or none as the value of the available permissions within the permissions key.

```
permissions:
    actions: read|write|none
    attestations: read|write|none
    checks: read|write|none
    contents: read|write|none
    deployments: read|write|none
    id—token: write|none
    issues: read|write|none
    models: read|none
    discussions: read|write|none
    packages: read|write|none
    pages: read|write|none
    pull—requests: read|write|none
    security—events: read|write|none
    statuses: read|write|none
```

If you specify the access for any of these permissions, all of those that are not specified are set to none.

You can use the following syntax to define one of read-all or write-all access for all of the available permissions:

```
permissions: read-all
permissions: write-all
```

You can use the following syntax to disable permissions for all of the available permissions:

```
permissions: {}
```

Changing the permissions in a forked repository \mathscr{D}

You can use the permissions key to add and remove read permissions for forked repositories, but typically you can't grant write access. The exception to this behavior is where an admin user

has selected the **Send write tokens to workflows from pull requests** option in the GitHub Actions settings. For more information, see Managing GitHub Actions settings for a repository.

Example: Setting the **GITHUB_TOKEN** permissions for one job in a workflow *P*

This example shows permissions being set for the GITHUB_TOKEN that will only apply to the job named stale. Write access is granted for the issues and pull-requests permissions. All other permissions will have no access.

```
jobs:
    stale:
        runs-on: ubuntu-latest

permissions:
        issues: write
        pull-requests: write

steps:
        - uses: actions/stale@v9
```

jobs.<job_id>.needs ∂

Use <code>jobs.<job_id>.needs</code> to identify any jobs that must complete successfully before this job will run. It can be a string or array of strings. If a job fails or is skipped, all jobs that need it are skipped unless the jobs use a conditional expression that causes the job to continue. If a run contains a series of jobs that need each other, a failure or skip applies to all jobs in the dependency chain from the point of failure or skip onwards. If you would like a job to run even if a job it is dependent on did not succeed, use the <code>always()</code> conditional expression in <code>jobs.</code> <code><job_id>.if</code>.

Example: Requiring successful dependent jobs 🔗

```
jobs:
    job1:
    job2:
     needs: job1
    job3:
     needs: [job1, job2]
```

In this example, job1 must complete successfully before job2 begins, and job3 waits for both job1 and job2 to complete.

The jobs in this example run sequentially:

- 1 job1
- 2 job2
- 3 job3

Example: Not requiring successful dependent jobs 🔗

```
jobs:
    job1:
    job2:
    needs: job1
    job3:
    if: ${{ always() }}
    needs: [job1, job2]
```

In this example, job3 uses the always() conditional expression so that it always runs after job1 and job2 have completed, regardless of whether they were successful. For more information, see Evaluate expressions in workflows and actions.

jobs.<job_id>.if ∂

You can use the <code>jobs.<job_id>.if</code> conditional to prevent a job from running unless a condition is met. You can use any supported context and expression to create a conditional. For more information on which contexts are supported in this key, see <code>Contexts</code> reference.

Note

The jobs.<job_id>.if condition is evaluated before jobs.<job_id>.strategy.matrix is applied.

When you use expressions in an if conditional, you can, optionally, omit the \${{}} expression syntax because GitHub Actions automatically evaluates the if conditional as an expression. However, this exception does not apply everywhere.

You must always use the \${{ }} expression syntax or escape with '', "", or () when the expression starts with !, since ! is reserved notation in YAML format. For example:

```
if: ${{ ! startsWith(github.ref, 'refs/tags/') }}
```

For more information, see Evaluate expressions in workflows and actions.

Example: Only run job for specific repository &

This example uses if to control when the production-deploy job can run. It will only run if the repository is named octo-repo-prod and is within the octo-org organization. Otherwise, the job will be marked as *skipped*.

```
name: example-workflow
on: [push]
jobs:
    production-deploy:
    if: github.repository == 'octo-org/octo-repo-prod'
    runs-on: ubuntu-latest
    steps:
        - uses: actions/checkout@v4
        - uses: actions/setup-node@v4
        with:
            node-version: '14'
        - run: npm install -g bats
```

jobs.<job_id>.runs-on ∂

Use jobs.<job_id>.runs-on to define the type of machine to run the job on.

- The destination machine can be either a <u>GitHub-hosted runner</u>, <u>larger runner</u>, or a <u>self-hosted</u> runner.
- You can target runners based on the labels assigned to them, or their group membership, or a combination of these.
- You can provide runs-on as:
 - A single string
 - A single variable containing a string
 - An array of strings, variables containing strings, or a combination of both

- A key: value pair using the group or labels keys
- If you specify an array of strings or variables, your workflow will execute on any runner that matches all of the specified runs—on values. For example, here the job will only run on a self-hosted runner that has the labels linux, x64, and gpu:

```
runs-on: [self-hosted, linux, x64, gpu]
```

For more information, see Choosing self-hosted runners.

• You can mix strings and variables in an array. For example:

```
on:
    workflow_dispatch:
    inputs:
        chosen-os:
        required: true
        type: choice
        options:
        - Ubuntu
        - macOS

jobs:
    test:
    runs-on: [self-hosted, "${{ inputs.chosen-os }}"]
    steps:
        - run: echo Hello world!
```

• If you would like to run your workflow on multiple machines, use jobs.<job_id>.strategy.

Note

Quotation marks are not required around simple strings like self-hosted, but they are required for expressions like "\${{ inputs.chosen-os }}".

Choosing GitHub-hosted runners *⊘*

If you use a GitHub-hosted runner, each job runs in a fresh instance of a runner image specified by runs-on.

The value for runs-on, when you are using a GitHub-hosted runner, is a runner label or the name of a runner group. The labels for the standard GitHub-hosted runners are shown in the following tables.

For more information, see GitHub-hosted runners.

Standard GitHub-hosted runners for public repositories *∂*

For public repositories, jobs using the workflow labels shown in the table below will run on virtual machines with the associated specifications. The use of these runners on public repositories is free and unlimited.

Virtual Machine	Processor (CPU)	Memory (RAM)	Storage (SSD)	Architecture	Workflow label
Linux	4	16 GB	14 GB	x64	ubuntu- latest , ubuntu-24.04 , ubuntu-22.04
Windows	4	16 GB	14 GB	x64	windows- latest , windows-2025 , windows-2022
Linux	4	16 GB	14 GB	arm64	ubuntu-24.04- arm , ubuntu- 22.04-arm
Windows	4	16 GB	14 GB	arm64	windows-11-
macOS	4	14 GB	14 GB	Intel	macos-13
macOS	3 (M1)	7 GB	14 GB	arm64	<pre>macos-latest , macos-14 , macos-15</pre>

Standard GitHub-hosted runners for private repositories *⊘*

For private repositories, jobs using the workflow labels shown in the table below will run on virtual machines with the associated specifications. These runners use your GitHub account's allotment of free minutes, and are then charged at the per minute rates. See Actions minute multiplier reference.

Virtual Machine	Processor (CPU)	Memory (RAM)	Storage (SSD)	Architecture	Workflow label
Linux	2	7 GB	14 GB	x64	ubuntu- latest , ubuntu-24.04 , ubuntu-22.04
Windows	2	7 GB	14 GB	x64	windows- latest , windows-2025 , windows-2022
macOS	4	14 GB	14 GB	Intel	macos-13
macOS	3 (M1)	7 GB	14 GB	arm64	<pre>macos-latest , macos-14 , macos-15</pre>

In addition to the standard GitHub-hosted runners, GitHub offers customers on GitHub Team and GitHub Enterprise Cloud plans a range of managed virtual machines with advanced features - for example, more cores and disk space, GPU-powered machines, and ARM-powered machines. For more information, see Larger runners.

Note

The —latest runner images are the latest stable images that GitHub provides, and might not be the most recent version of the operating system available from the operating system vendor.

Beta and Deprecated Images are provided "as-is", "with all faults" and "as available" and are excluded from the service level agreement and warranty. Beta Images may not be covered by customer support.

Example: Specifying an operating system \mathscr{D}

runs-on: ubuntu-latest

For more information, see GitHub-hosted runners.

Choosing self-hosted runners *∂*

To specify a self-hosted runner for your job, configure runs-on in your workflow file with self-hosted runner labels.

Self-hosted runners may have the self-hosted label. When setting up a self-hosted runner, by default we will include the label self-hosted. You may pass in the --no-default-labels flag to prevent the self-hosted label from being applied. Labels can be used to create targeting options for runners, such as operating system or architecture, we recommend providing an array of labels that begins with self-hosted (this must be listed first) and then includes additional labels as needed. When you specify an array of labels, jobs will be queued on runners that have all the labels that you specify.

Note that Actions Runner Controller does not support multiple labels and does not support the self-hosted label.

Example: Using labels for runner selection \mathscr{O}

```
runs-on: [self-hosted, linux]
```

For more information, see Self-hosted runners and Using self-hosted runners in a workflow.

Choosing runners in a group *∂*

You can use runs-on to target runner groups, so that the job will execute on any runner that is a member of that group. For more granular control, you can also combine runner groups with labels.

Runner groups can only have larger runners or self-hosted runners as members.

Example: Using groups to control where jobs are run \mathscr{S}

In this example, Ubuntu runners have been added to a group called ubuntu-runners. The runson key sends the job to any available runner in the ubuntu-runners group:

```
name: learn-github-actions
on: [push]
jobs:
   check-bats-version:
    runs-on:
      group: ubuntu-runners
```

```
steps:
    - uses: actions/checkout@v4
    - uses: actions/setup-node@v4
    with:
        node-version: '14'
    - run: npm install -g bats
    - run: bats -v
```

Example: Combining groups and labels &

When you combine groups and labels, the runner must meet both requirements to be eligible to run the job.

In this example, a runner group called ubuntu-runners is populated with Ubuntu runners, which have also been assigned the label ubuntu-20.04-16core. The runs-on key combines group and labels so that the job is routed to any available runner within the group that also has a matching label:

```
name: learn-github-actions
on: [push]
jobs:
    check-bats-version:
    runs-on:
        group: ubuntu-runners
        labels: ubuntu-20.04-16core
    steps:
        - uses: actions/checkout@v4
        - uses: actions/setup-node@v4
        with:
            node-version: '14'
        - run: npm install -g bats
        - run: bats -v
```

jobs.<job_id>.environment ∂

Use jobs.<job_id>.environment to define the environment that the job references.

You can provide the environment as only the environment name, or as an environment object with the name and url. The URL maps to environment_url in the deployments API. For more information about the deployments API, see REST API endpoints for repositories.



All deployment protection rules must pass before a job referencing the environment is sent to a runner. For more information, see Managing environments for deployment.

Example: Using a single environment name *∂*

```
environment: staging_environment
```

Example: Using environment name and URL ∂

```
environment:
  name: production_environment
  url: https://github.com
```

The value of url can be an expression. Allowed expression contexts: github, inputs, vars, needs, strategy, matrix, job, runner, env, and steps. For more information about expressions, see Evaluate expressions in workflows and actions.

Example: Using output as URL *∂*

```
environment:
  name: production_environment
  url: ${{ steps.step_id.outputs.url_output }}
```

The value of name can be an expression. Allowed expression contexts: github, inputs, vars, needs, strategy, and matrix. For more information about expressions, see Evaluate expressions in workflows and actions.

Example: Using an expression as environment name *∂*

```
environment:
  name: ${{ github.ref_name }}
```

jobs.<job_id>.concurrency ∂

You can use <code>jobs.<job_id>.concurrency</code> to ensure that only a single job or workflow using the same concurrency group will run at a time. A concurrency group can be any string or expression. Allowed expression contexts: <code>github</code>, <code>inputs</code>, <code>vars</code>, <code>needs</code>, <code>strategy</code>, and <code>matrix</code>. For more information about expressions, see Evaluate expressions in workflows and actions.

You can also specify concurrency at the workflow level. For more information, see concurrency.

This means that there can be at most one running and one pending job in a concurrency group at any time. When a concurrent job or workflow is queued, if another job or workflow using the same concurrency group in the repository is in progress, the queued job or workflow will be pending. Any existing pending job or workflow in the same concurrency group, if it exists, will be canceled and the new queued job or workflow will take its place.

To also cancel any currently running job or workflow in the same concurrency group, specify cancel—in—progress: true. To conditionally cancel currently running jobs or workflows in the same concurrency group, you can specify cancel—in—progress as an expression with any of the allowed expression contexts.

(i) Note

- The concurrency group name is case insensitive. For example, prod and Prod will be treated as the same concurrency group.
- Ordering is not guaranteed for jobs or workflow runs using concurrency groups. Jobs or workflow runs in the same concurrency group are handled in an arbitrary order.

Example: Using concurrency and the default behavior &

The default behavior of GitHub Actions is to allow multiple jobs or workflow runs to run concurrently. The concurrency keyword allows you to control the concurrency of workflow runs.

For example, you can use the concurrency keyword immediately after where trigger conditions are defined to limit the concurrency of entire workflow runs for a specific branch:

```
on:
    push:
        branches:
        - main

concurrency:
    group: ${{ github.workflow }}-${{ github.ref }}
    cancel-in-progress: true
```

You can also limit the concurrency of jobs within a workflow by using the concurrency keyword at the job level:

```
on:
    push:
    branches:
    - main

jobs:
    job-1:
    runs-on: ubuntu-latest
    concurrency:
        group: example-group
        cancel-in-progress: true
```

Example: Concurrency groups $\mathscr O$

Concurrency groups provide a way to manage and limit the execution of workflow runs or jobs that share the same concurrency key.

The concurrency key is used to group workflows or jobs together into a concurrency group. When you define a concurrency key, GitHub Actions ensures that only one workflow or job with that key runs at any given time. If a new workflow run or job starts with the same concurrency key, GitHub Actions will cancel any workflow or job already running with that key. The concurrency key can be a hard-coded string, or it can be a dynamic expression that includes context variables.

It is possible to define concurrency conditions in your workflow so that the workflow or job is part of a concurrency group.

This means that when a workflow run or job starts, GitHub will cancel any workflow runs or jobs that are already in progress in the same concurrency group. This is useful in scenarios where you want to prevent parallel runs for a certain set of a workflows or jobs, such as the ones used for deployments to a staging environment, in order to prevent actions that could cause conflicts or consume more resources than necessary.

In this example, job-1 is part of a concurrency group named staging_environment. This means that if a new run of job-1 is triggered, any runs of the same job in the staging_environment concurrency group that are already in progress will be cancelled.

```
jobs:
job-1:
```

```
runs-on: ubuntu-latest
concurrency:
group: staging_environment
cancel-in-progress: true
```

Alternatively, using a dynamic expression such as concurrency: ci-\${{ github.ref }} in your workflow means that the workflow or job would be part of a concurrency group named cifollowed by the reference of the branch or tag that triggered the workflow. In this example, if a new commit is pushed to the main branch while a previous run is still in progress, the previous run will be cancelled and the new one will start:

```
on:
   push:
    branches:
    - main

concurrency:
   group: ci-${{ github.ref }}
   cancel-in-progress: true
```

Example: Using concurrency to cancel any in-progress job or run &

To use concurrency to cancel any in-progress job or run in GitHub Actions, you can use the concurrency key with the cancel-in-progress option set to true:

```
concurrency:
  group: ${{ github.ref }}
  cancel-in-progress: true
```

Note that in this example, without defining a particular concurrency group, GitHub Actions will cancel *any* in-progress run of the job or workflow.

Example: Using a fallback value *∂*

If you build the group name with a property that is only defined for specific events, you can use a fallback value. For example, <code>github.head_ref</code> is only defined on <code>pull_request</code> events. If your workflow responds to other events in addition to <code>pull_request</code> events, you will need to provide a fallback to avoid a syntax error. The following concurrency group cancels in-progress jobs or runs on <code>pull_request</code> events only; if <code>github.head_ref</code> is undefined, the concurrency group will fallback to the run ID, which is guaranteed to be both unique and defined for the run.

```
concurrency:
  group: ${{ github.head_ref || github.run_id }}
  cancel-in-progress: true
```

Example: Only cancel in-progress jobs or runs for the current workflow &

If you have multiple workflows in the same repository, concurrency group names must be unique across workflows to avoid canceling in-progress jobs or runs from other workflows. Otherwise, any previously in-progress or pending job will be canceled, regardless of the workflow.

To only cancel in-progress runs of the same workflow, you can use the github.workflow property to build the concurrency group:

```
concurrency:
  group: ${{ github.workflow }}-${{ github.ref }}
  cancel-in-progress: true
```

Example: Only cancel in-progress jobs on specific branches *∂*

If you would like to cancel in-progress jobs on certain branches but not on others, you can use conditional expressions with <code>cancel-in-progress</code>. For example, you can do this if you would like to cancel in-progress jobs on development branches but not on release branches.

To only cancel in-progress runs of the same workflow when not running on a release branch, you can set cancel-in-progress to an expression similar to the following:

```
concurrency:
  group: ${{ github.workflow }}-${{ github.ref }}
  cancel-in-progress: ${{ !contains(github.ref, 'release/')}}
```

In this example, multiple pushes to a release/1.2.3 branch would not cancel in-progress runs. Pushes to another branch, such as main, would cancel in-progress runs.

jobs.<job_id>.outputs ∂

You can use <code>jobs.<job_id>.outputs</code> to create a map of outputs for a job. Job outputs are available to all downstream jobs that depend on this job. For more information on defining job dependencies, see <code>jobs.<job_id>.needs</code>.

Outputs can be a maximum of 1 MB per job. The total of all outputs in a workflow run can be a maximum of 50 MB. Size is approximated based on UTF-16 encoding.

Job outputs containing expressions are evaluated on the runner at the end of each job. Outputs containing secrets are redacted on the runner and not sent to GitHub Actions.

If an output is skipped because it may contain a secret, you will see the following warning message: "Skip output {output.Key} since it may contain secret." For more information on how to handle secrets, please refer to the Example: Masking and passing a secret between jobs or workflows.

To use job outputs in a dependent job, you can use the needs context. For more information, see Contexts reference.

Example: Defining outputs for a job *⊘*

```
jobs:
  job1:
    runs-on: ubuntu-latest
   # Map a step output to a job output
    outputs:
      output1: ${{ steps.step1.outputs.test }}
      output2: ${{ steps.step2.outputs.test }}
    steps:
      - id: step1
        run: echo "test=hello" >> "$GITHUB_OUTPUT"
     - id: step2
        run: echo "test=world" >> "$GITHUB_OUTPUT"
  job2:
    runs-on: ubuntu-latest
    needs: job1
    steps:
      - env:
          OUTPUT1: ${{needs.job1.outputs.output1}}
          OUTPUT2: ${{needs.job1.outputs.output2}}
        run: echo "$0UTPUT1 $0UTPUT2"
```

Using Job Outputs in a Matrix Job *∂*

Matrices can be used to generate multiple outputs of different names. When using a matrix, job outputs will be combined from all jobs inside the matrix.

```
jobs:
  job1:
    runs-on: ubuntu-latest
    outputs:
      output_1: ${{ steps.gen_output.outputs.output_1 }}
      output 2: ${{ steps.gen output.outputs.output 2 }}
      output_3: ${{ steps.gen_output.outputs.output_3 }}
    strategy:
     matrix:
        version: [1, 2, 3]
    steps:
      - name: Generate output
        id: gen_output
        run:
          version="${{ matrix.version }}"
          echo "output_${version}=${version}" >> "$GITHUB_OUTPUT"
  job2:
    runs-on: ubuntu-latest
    needs: [job1]
    steps:
      # Will show
     # {
          "output_1": "1",
          "output 2": "2",
         "output 3": "3"
      - run: echo '${{ toJSON(needs.job1.outputs) }}'
```

Warning

Actions does not guarantee the order that matrix jobs will run in. Ensure that the output name is unique, otherwise the last matrix job that runs will override the output value.

jobs.<job_id>.env ∂

A map of variables that are available to all steps in the job. You can set variables for the entire workflow or an individual step. For more information, see env and jobs. <a href="

When more than one environment variable is defined with the same name, GitHub uses the most specific variable. For example, an environment variable defined in a step will override job and workflow environment variables with the same name, while the step executes. An environment variable defined for a job will override a workflow variable with the same name, while the job executes.

Example of jobs.<job_id>.env ∂

```
jobs:
   job1:
   env:
    FIRST_NAME: Mona
```

jobs.<job_id>.defaults ∂

Use <code>jobs.<job_id>.defaults</code> to create a map of default settings that will apply to all steps in the job. You can also set default settings for the entire workflow. For more information, see <code>defaults</code>.

When more than one default setting is defined with the same name, GitHub uses the most specific default setting. For example, a default setting defined in a job will override a default setting that has the same name defined in a workflow.

jobs.<job_id>.defaults.run ∂

Use jobs.<job_id>.defaults.run to provide default shell and working-directory to all run steps in the job.

You can provide default shell and working-directory options for all <u>run</u> steps in a job. You can also set default settings for run for the entire workflow. For more information, see <u>defaults.run</u>.

These can be overridden at the jobs.<job_id>.defaults.run and jobs.<job_id>.steps[*].run levels.

When more than one default setting is defined with the same name, GitHub uses the most specific default setting. For example, a default setting defined in a job will override a default setting that has the same name defined in a workflow.

jobs.<job_id>.defaults.run.shell ₽

Use shell to define the shell for a step. This keyword can reference several contexts. For more information, see Contexts.

Supported platform	shell parameter	Description	Command run internally
Linux / macOS	unspecified	The default shell on non-Windows platforms. Note that this runs a different command to when bash is specified explicitly. If bash is not found in the path, this is treated as sh.	bash -e {0}
All	bash	The default shell on non-Windows platforms with a fallback to sh. When specifying a bash shell on Windows, the bash shell included with Git for Windows is used.	<pre>bashnoprofile norc -eo pipefail {0}</pre>
All	pwsh	The PowerShell Core. GitHub appends the extension .ps1 to your script name.	<pre>pwsh -command ". '{0}'"</pre>
All	python	Executes the python command.	python {0}
Linux / macOS	sh	The fallback behavior for non-Windows platforms if no shell is provided and bash is not found in the path.	sh -e {0}
Windows	cmd	GitHub appends the extension .cmd to your script name and substitutes for {0}.	%ComSpec% /D /E:ON /V:OFF /S /C "CALL " {0}"".
Windows	pwsh	This is the default shell used on Windows. The PowerShell Core. GitHub appends the extension .ps1 to your script	<pre>pwsh -command ". '{0}'".</pre>

Supported platform	shell parameter	Description	Command run internally
		name. If your self- hosted Windows runner does not have PowerShell Core installed, then PowerShell Desktop is used instead.	
Windows	powershell	The PowerShell Desktop. GitHub appends the extension .ps1 to your script name.	<pre>powershell -command ". '{0}'".</pre>

When more than one default setting is defined with the same name, GitHub uses the most specific default setting. For example, a default setting defined in a job will override a default setting that has the same name defined in a workflow.

jobs.<job_id>.defaults.run.working-directory ∂

Use working-directory to define the working directory for the shell for a step. This keyword can reference several contexts. For more information, see Contexts.

Ω Tip

Ensure the working-directory you assign exists on the runner before you run your shell in it. When more than one default setting is defined with the same name, GitHub uses the most specific default setting. For example, a default setting defined in a job will override a default setting that has the same name defined in a workflow.

Example: Setting default run step options for a job 🔗

```
jobs:
    job1:
    runs-on: ubuntu-latest
    defaults:
    run:
       shell: bash
       working-directory: ./scripts
```

jobs.<job_id>.steps ∂

A job contains a sequence of tasks called steps. Steps can run commands, run setup tasks, or run an action in your repository, a public repository, or an action published in a Docker registry. Not all steps run actions, but all actions run as a step. Each step runs in its own process in the runner environment and has access to the workspace and filesystem. Because steps run in their own process, changes to environment variables are not preserved between steps. GitHub provides built-in steps to set up and complete a job.

GitHub only displays the first 1,000 checks, however, you can run an unlimited number of steps as long as you are within the workflow usage limits. For more information, see <u>Billing and usage</u> for GitHub-hosted runners and <u>Actions limits</u> for self-hosted runner usage limits.

Example of jobs.<job_id>.steps ∂

```
name: Greeting from Mona
on: push

jobs:
    my-job:
    name: My Job
    runs-on: ubuntu-latest
    steps:
        - name: Print a greeting
        env:
            MY_VAR: Hi there! My name is
            FIRST_NAME: Mona
            MIDDLE_NAME: The
            LAST_NAME: Octocat
    run: |
            echo $MY_VAR $FIRST_NAME $MIDDLE_NAME $LAST_NAME.
```

jobs.<job_id>.steps[*].id ₽

A unique identifier for the step. You can use the id to reference the step in contexts. For more information, see Contexts reference.

jobs.<job_id>.steps[*].if ♂

You can use the if conditional to prevent a step from running unless a condition is met. You can use any supported context and expression to create a conditional. For more information on which contexts are supported in this key, see Contexts reference.

When you use expressions in an if conditional, you can, optionally, omit the \${{}} expression syntax because GitHub Actions automatically evaluates the if conditional as an expression. However, this exception does not apply everywhere.

You must always use the \${{ }} expression syntax or escape with '', "", or () when the expression starts with !, since ! is reserved notation in YAML format. For example:

```
if: ${{ ! startsWith(github.ref, 'refs/tags/') }}
```

For more information, see Evaluate expressions in workflows and actions.

Example: Using contexts &

This step only runs when the event type is a <code>pull_request</code> and the event action is <code>unassigned</code>.

```
steps:
    - name: My first step
    if: ${{ github.event_name == 'pull_request' && github.event.action == 'unassigned' }}
    run: echo This event is a pull request that had an assignee removed.
```

Example: Using status check functions ∂

The my backup step only runs when the previous step of a job fails. For more information, see Evaluate expressions in workflows and actions.

```
steps:
    - name: My first step
    uses: octo-org/action-name@main
    - name: My backup step
    if: ${{ failure() }}
    uses: actions/heroku@1.0.0
```

Example: Using secrets *⊘*

Secrets cannot be directly referenced in if: conditionals. Instead, consider setting secrets as job-level environment variables, then referencing the environment variables to conditionally run steps in the job.

If a secret has not been set, the return value of an expression referencing the secret (such as \${{ secrets.SuperSecret }} in the example) will be an empty string.

```
name: Run a step if a secret has been set
on: push
jobs:
    my-jobname:
    runs-on: ubuntu-latest
    env:
        super_secret: ${{ secrets.SuperSecret }}
    steps:
        - if: ${{ env.super_secret != '' }}
        run: echo 'This step will only run if the secret has a value set.'
        - if: ${{ env.super_secret == '' }}
        run: echo 'This step will only run if the secret does not have a value set.'
```

For more information, see Contexts reference and Using secrets in GitHub Actions.

jobs.<job_id>.steps[*].name ₽

A name for your step to display on GitHub.

jobs.<job_id>.steps[*].uses ₽

Selects an action to run as part of a step in your job. An action is a reusable unit of code. You can use an action defined in the same repository as the workflow, a public repository, or in a <u>published</u> Docker container image.

We strongly recommend that you include the version of the action you are using by specifying a Git ref, SHA, or Docker tag. If you don't specify a version, it could break your workflows or cause unexpected behavior when the action owner publishes an update.

- Using the commit SHA of a released action version is the safest for stability and security.
- If the action publishes major version tags, you should expect to receive critical fixes and security patches while still retaining compatibility. Note that this behavior is at the discretion of the action's author.

• Using the default branch of an action may be convenient, but if someone releases a new major version with a breaking change, your workflow could break.

Some actions require inputs that you must set using the with keyword. Review the action's README file to determine the inputs required.

Actions are either JavaScript files or Docker containers. If the action you're using is a Docker container you must run the job in a Linux environment. For more details, see runs-on.

Example: Using versioned actions *∂*

```
steps:
    # Reference a specific commit
    - uses: actions/checkout@8f4b7f84864484a7bf31766abe9204da3cbe65b3
# Reference the major version of a release
    - uses: actions/checkout@v4
# Reference a specific version
    - uses: actions/checkout@v4.2.0
# Reference a branch
    - uses: actions/checkout@main
```

Example: Using a public action *∂*

```
{owner}/{repo}@{ref}
```

You can specify a branch, ref, or SHA in a public GitHub repository.

```
jobs:
    my_first_job:
    steps:
        - name: My first step
        # Uses the default branch of a public repository
        uses: actions/heroku@main
        - name: My second step
        # Uses a specific version tag of a public repository
        uses: actions/aws@v2.0.1
```

Example: Using a public action in a subdirectory *∂*

{owner}/{repo}/{path}@{ref}

A subdirectory in a public GitHub repository at a specific branch, ref, or SHA.

```
jobs:
    my_first_job:
    steps:
    - name: My first step
    uses: actions/aws/ec2@main
```

Example: Using an action in the same repository as the workflow *?*

```
./path/to/dir
```

The path to the directory that contains the action in your workflow's repository. You must check out your repository before using the action.

Example repository file structure:

The path is relative (./) to the default working directory (github.workspace , \$GITHUB_WORKSPACE). If the action checks out the repository to a location different than the workflow, the relative path used for local actions must be updated.

Example workflow file:

```
jobs:
    my_first_job:
    runs-on: ubuntu-latest
    steps:
        # This step checks out a copy of your repository.
        - name: My first step - check out repository
        uses: actions/checkout@v4
        # This step references the directory that contains the action.
        - name: Use local hello-world-action
        uses: ./.github/actions/hello-world-action
```

Example: Using a Docker Hub action *∂*

```
docker://{image}:{tag}
```

A Docker image published on Docker Hub.

```
jobs:
    my_first_job:
    steps:
        - name: My first step
        uses: docker://alpine:3.8
```

Example: Using the GitHub Packages Container registry $\mathscr D$

```
docker://{host}/{image}:{tag}
```

A public Docker image in the GitHub Packages Container registry.

```
jobs:
    my_first_job:
    steps:
    - name: My first step
    uses: docker://ghcr.io/OWNER/IMAGE_NAME
```

Example: Using a Docker public registry action *∂*

```
docker://{host}/{image}:{tag}
```

A Docker image in a public registry. This example uses the Google Container Registry at gcr.io.

```
jobs:
    my_first_job:
    steps:
        - name: My first step
        uses: docker://gcr.io/cloud-builders/gradle
```

Example: Using an action inside a different private repository than the workflow \mathscr{P}

Your workflow must checkout the private repository and reference the action locally. Generate a personal access token and add the token as a secret. For more information, see <u>Managing your</u> personal access tokens and Using secrets in GitHub Actions.

Replace PERSONAL_ACCESS_TOKEN in the example with the name of your secret.

```
jobs:
    my_first_job:
    steps:
        - name: Check out repository
        uses: actions/checkout@v4
        with:
            repository: octocat/my-private-repo
            ref: v1.0
            token: ${{ secrets.PERSONAL_ACCESS_TOKEN }}
            path: ./.github/actions/my-private-repo
            - name: Run my action
            uses: ./.github/actions/my-private-repo/my-action
```

Alternatively, use a GitHub App instead of a personal access token in order to ensure your workflow continues to run even if the personal access token owner leaves. For more information, see Making authenticated API requests with a GitHub App in a GitHub Actions workflow.

jobs.<job_id>.steps[*].run ♂

Runs command-line programs that do not exceed 21,000 characters using the operating system's shell. If you do not provide a name, the step name will default to the text specified in the run command.

Commands run using non-login shells by default. You can choose a different shell and customize the shell used to run commands. For more information, see jobs.<job_id>.steps[*].shell.

Each run keyword represents a new process and shell in the runner environment. When you provide multi-line commands, each line runs in the same shell. For example:

• A single-line command:

```
- name: Install Dependencies
run: npm install
```

• A multi-line command:

```
- name: Clean install dependencies and build
run: |
   npm ci
   npm run build
```

jobs.<job_id>.steps[*].working-directory ♂

Using the working-directory keyword, you can specify the working directory of where to run the command.

```
- name: Clean temp directory
run: rm -rf *
working-directory: ./temp
```

Alternatively, you can specify a default working directory for all run steps in a job, or for all run steps in the entire workflow. For more information, see defaults.run.working-directory and <a href="jobs.<jobs.defaults.run.working-directory">jobs.<jobs.defaults.run.working-directory.

You can also use a run step to run a script. For more information, see Adding scripts to your workflow.

jobs.<job_id>.steps[*].shell ♂

You can override the default shell settings in the runner's operating system and the job's default using the shell keyword. You can use built-in shell keywords, or you can define a custom set of shell options. The shell command that is run internally executes a temporary file that contains the commands specified in the run keyword.

Supported platform	shell parameter	Description	Command run internally
Linux / macOS	unspecified	The default shell on non-Windows platforms. Note that this runs a different command to when bash is specified explicitly. If bash is not found in the path, this is treated as sh.	bash -e {0}
All	bash	The default shell on non-Windows platforms with a fallback to sh. When specifying a bash shell on Windows, the bash shell included with Git for Windows is used.	<pre>bashnoprofile norc -eo pipefail {0}</pre>
All	pwsh	The PowerShell Core. GitHub appends the extension .ps1 to your script name.	<pre>pwsh -command ". '{0}'"</pre>
All	python	Executes the python command.	python {0}
Linux / macOS	sh	The fallback behavior for non-Windows platforms if no shell is provided and bash is not found in the path.	sh -e {0}
Windows	cmd	GitHub appends the extension .cmd to your script name and substitutes for {0}.	%ComSpec% /D /E:ON /V:OFF /S /C "CALL " {0}"".
Windows	pwsh	This is the default shell used on Windows. The PowerShell Core. GitHub appends the extension .ps1 to your script	<pre>pwsh -command ". '{0}'".</pre>

Supported platform	shell parameter	Description	Command run internally
		name. If your self- hosted Windows runner does not have PowerShell Core installed, then PowerShell Desktop is used instead.	
Windows	powershell	The PowerShell Desktop. GitHub appends the extension .ps1 to your script name.	powershell -command ". '{0}'".

Alternatively, you can specify a default shell for all run steps in a job, or for all run steps in the entire workflow. For more information, see defaults.run.shell and jobs.defaults.run.shell.

Example: Running a command using Bash ∂

steps:

- name: Display the path

shell: bash
run: echo \$PATH

Example: Running a command using Windows cmd ?

steps:

- name: Display the path

shell: cmd

run: echo %PATH%

Example: Running a command using PowerShell Core &

steps:

- name: Display the path

```
shell: pwsh
run: echo ${env:PATH}
```

Example: Using PowerShell Desktop to run a command *?*

```
steps:
    - name: Display the path
    shell: powershell
    run: echo ${env:PATH}
```

Example: Running an inline Python script *⊘*

```
steps:
   - name: Display the path
   shell: python
   run: |
    import os
    print(os.environ['PATH'])
```

Custom shell &

You can set the shell value to a template string using command [options] {0} [more_options] . GitHub interprets the first whitespace-delimited word of the string as the command, and inserts the file name for the temporary script at {0}.

For example:

```
steps:
   - name: Display the environment variables and their values
   shell: perl {0}
   run: |
      print %ENV
```

The command used, perl in this example, must be installed on the runner.

For information about the software included on GitHub-hosted runners, see <u>GitHub-hosted</u> runners.

Exit codes and error action preference *⊘*

For built-in shell keywords, we provide the following defaults that are executed by GitHub-hosted runners. You should use these guidelines when running shell scripts.

bash / sh :

- By default, fail-fast behavior is enforced using set -e for both sh and bash. When shell: bash is specified, -o pipefail is also applied to enforce early exit from pipelines that generate a non-zero exit status.
- You can take full control over shell parameters by providing a template string to the shell options. For example, bash {0}.
- sh -like shells exit with the exit code of the last command executed in a script, which is also the default behavior for actions. The runner will report the status of the step as fail/succeed based on this exit code.

powershell / pwsh

- Fail-fast behavior when possible. For pwsh and powershell built-in shell, we will prepend \$ErrorActionPreference = 'stop' to script contents.
- We append if ((Test-Path -LiteralPath variable:\LASTEXITCODE)) { exit
 \$LASTEXITCODE } to powershell scripts so action statuses reflect the script's last exit code.
- Users can always opt out by not using the built-in shell, and providing a custom shell
 option like: pwsh -File {0}, or powershell -Command "& '{0}'", depending on need.

cmd

- There doesn't seem to be a way to fully opt into fail-fast behavior other than writing your script to check each error code and respond accordingly. Because we can't actually provide that behavior by default, you need to write this behavior into your script.
- o cmd.exe will exit with the error level of the last program it executed, and it will return the error code to the runner. This behavior is internally consistent with the previous sh and pwsh default behavior and is the cmd.exe default, so this behavior remains intact.

jobs.<job_id>.steps[*].with ∂

A map of the input parameters defined by the action. Each input parameter is a key/value pair. Input parameters are set as environment variables. The variable is prefixed with INPUT_ and converted to upper case.

Input parameters defined for a Docker container must use args. For more information, see jobs.<job_id>.steps[*].with.args.

Example of jobs.<job_id>.steps[*].with ∂

Defines the three input parameters (first_name, middle_name, and last_name) defined by the hello_world action. These input variables will be accessible to the hello-world action as INPUT_FIRST_NAME, INPUT_MIDDLE_NAME, and INPUT_LAST_NAME environment variables.

```
jobs:
    my_first_job:
    steps:
        - name: My first step
        uses: actions/hello_world@main
        with:
        first_name: Mona
        middle_name: The
        last_name: Octocat
```

jobs.<job_id>.steps[*].with.args ♂

A string that defines the inputs for a Docker container. GitHub passes the args to the container's ENTRYPOINT when the container starts up. An array of strings is not supported by this parameter. A single argument that includes spaces should be surrounded by double quotes "".

Example of jobs.<job_id>.steps[*].with.args ∂

```
steps:
   - name: Explain why this job ran
   uses: octo-org/action-name@main
   with:
      entrypoint: /bin/echo
      args: The ${{ github.event_name }} event triggered this step.
```

The args are used in place of the CMD instruction in a Dockerfile. If you use CMD in your Dockerfile, use the guidelines ordered by preference:

- 1 Document required arguments in the action's README and omit them from the CMD instruction.
- 2 Use defaults that allow using the action without specifying any args.
- 3 If the action exposes a —help flag, or something similar, use that as the default to make your action self-documenting.

jobs.<job_id>.steps[*].with.entrypoint ♂

Overrides the Docker ENTRYPOINT in the Dockerfile, or sets it if one wasn't already specified. Unlike the Docker ENTRYPOINT instruction which has a shell and exec form, entrypoint keyword accepts only a single string defining the executable to be run.

Example of jobs.<job_id>.steps[*].with.entrypoint ∂

steps:

- name: Run a custom command

uses: octo-org/action-name@main

with:

entrypoint: /a/different/executable

The entrypoint keyword is meant to be used with Docker container actions, but you can also use it with JavaScript actions that don't define any inputs.

jobs.<job_id>.steps[*].env ♂

Sets variables for steps to use in the runner environment. You can also set variables for the entire workflow or a job. For more information, see env and <a href="mailto:jobs. <a href="mailto:jobs.

When more than one environment variable is defined with the same name, GitHub uses the most specific variable. For example, an environment variable defined in a step will override job and workflow environment variables with the same name, while the step executes. An environment variable defined for a job will override a workflow variable with the same name, while the job executes.

Public actions may specify expected variables in the README file. If you are setting a secret or sensitive value, such as a password or token, you must set secrets using the secrets context. For more information, see Contexts reference.

Example of jobs.<job_id>.steps[*].env ∂

```
steps:
    - name: My first action
    env:
        GITHUB_TOKEN: ${{ secrets.GITHUB_TOKEN }}
        FIRST_NAME: Mona
        LAST_NAME: Octocat
```

jobs.<job_id>.steps[*].continue-on-error ∂

Prevents a job from failing when a step fails. Set to true to allow a job to pass when this step fails.

jobs.<job_id>.steps[*].timeout-minutes ∂

The maximum number of minutes to run the step before killing the process.

Fractional values are not supported. timeout-minutes must be a positive integer.

jobs.<job_id>.timeout-minutes ≥

The maximum number of minutes to let a job run before GitHub automatically cancels it. Default: 360

If the timeout exceeds the job execution time limit for the runner, the job will be canceled when the execution time limit is met instead. For more information about job execution time limits, see Billing and usage for GitHub-hosted runners and Actions limits for self-hosted runner usage limits.

Note

The GITHUB_TOKEN expires when a job finishes or after a maximum of 24 hours. For self-hosted runners, the token may be the limiting factor if the job timeout is greater than 24 hours. For more information on the GITHUB_TOKEN, see Use GITHUB_TOKEN for authentication in workflows.

jobs.<job_id>.strategy ♂

Use <code>jobs.<job_id>.strategy</code> to use a matrix strategy for your jobs. A matrix strategy lets you use variables in a single job definition to automatically create multiple job runs that are based on the combinations of the variables. For example, you can use a matrix strategy to test your code in multiple versions of a language or on multiple operating systems. For more information, see Running variations of jobs in a workflow.

jobs.<job_id>.strategy.matrix ♂

Use jobs.<job_id>.strategy.matrix to define a matrix of different job configurations. For more information, see Running variations of jobs in a workflow.

A matrix will generate a maximum of 256 jobs per workflow run. This limit applies to both GitHubhosted and self-hosted runners.

The variables that you define become properties in the matrix context, and you can reference the property in other areas of your workflow file. In this example, you can use matrix.version and matrix.os to access the current value of version and os that the job is using. For more information, see Contexts reference.

By default, GitHub will maximize the number of jobs run in parallel depending on runner availability. The order of the variables in the matrix determines the order in which the jobs are created. The first variable you define will be the first job that is created in your workflow run.

Using a single-dimension matrix *∂*

The following workflow defines the variable version with the values [10, 12, 14]. The workflow will run three jobs, one for each value in the variable. Each job will access the version value through the matrix.version context and pass the value as node-version to the actions/setup-node action.

```
jobs:
    example_matrix:
    strategy:
        matrix:
        version: [10, 12, 14]
    steps:
        - uses: actions/setup-node@v4
```

```
with:
  node-version: ${{ matrix.version }}
```

Using a multi-dimensional matrix *∂*

Specify multiple variables to create a multi-dimensional matrix. A job will run for each possible combination of the variables.

For example, the following workflow specifies two variables:

- Two operating systems specified in the os variable
- Three Node.js versions specified in the version variable

The workflow will run six jobs, one for each combination of the os and version variables. Each job will set the runs-on value to the current os value and will pass the current version value to the actions/setup-node action.

```
jobs:
    example_matrix:
        strategy:
        matrix:
        os: [ubuntu-22.04, ubuntu-20.04]
        version: [10, 12, 14]
    runs-on: ${{ matrix.os }}
    steps:
    - uses: actions/setup-node@v4
        with:
        node-version: ${{ matrix.version }}
```

A variable configuration in a matrix can be an array of object s. For example, the following matrix produces 4 jobs with corresponding contexts.

```
matrix:
    os:
        - ubuntu-latest
        - macos-latest
node:
        - version: 14
        - version: 20
        env: NODE_OPTIONS=--openssl-legacy-provider
```

Each job in the matrix will have its own combination of os and node values, as shown below.

```
- matrix.os: ubuntu-latest
  matrix.node.version: 14
- matrix.node.version: 20
  matrix.node.env: NODE_OPTIONS=--openssl-legacy-provider
- matrix.os: macos-latest
  matrix.node.version: 14
- matrix.os: macos-latest
  matrix.node.version: 20
  matrix.node.version: 20
  matrix.node.env: NODE_OPTIONS=--openssl-legacy-provider
```

jobs.<job_id>.strategy.matrix.include ♂

For each object in the include list, the key:value pairs in the object will be added to each of the matrix combinations if none of the key:value pairs overwrite any of the original matrix values. If the object cannot be added to any of the matrix combinations, a new matrix combination will be created instead. Note that the original matrix values will not be overwritten, but added matrix values can be overwritten.

Example: Expanding configurations *⊘*

For example, the following workflow will run four jobs, one for each combination of os and node. When the job for the os value of windows—latest and node value of 16 runs, an additional variable called npm with the value of 6 will be included in the job.

```
jobs:
 example matrix:
    strategy:
     matrix:
        os: [windows-latest, ubuntu-latest]
        node: [14, 16]
        include:
          - os: windows-latest
            node: 16
            npm: 6
    runs-on: ${{ matrix.os }}
    steps:
      - uses: actions/setup-node@v4
       with:
          node-version: ${{ matrix.node }}
     - if: ${{ matrix.npm }}
```

```
run: npm install -g npm@${{ matrix.npm }}
- run: npm --version
```

Example: Adding configurations \mathscr{D}

For example, this matrix will run 10 jobs, one for each combination of os and version in the matrix, plus a job for the os value of windows—latest and version value of 17.

If you don't specify any matrix variables, all configurations under include will run. For example, the following workflow would run two jobs, one for each include entry. This lets you take advantage of the matrix strategy without having a fully populated matrix.

```
jobs:
  includes_only:
    runs-on: ubuntu-latest
    strategy:
    matrix:
       include:
          - site: "production"
            datacenter: "site-a"
            - site: "staging"
            datacenter: "site-b"
```

jobs.<job_id>.strategy.matrix.exclude @

An excluded configuration only has to be a partial match for it to be excluded.

All include combinations are processed after exclude. This allows you to use include to add back combinations that were previously excluded.

jobs.<job_id>.strategy.fail-fast ♂

jobs.<job_id>.strategy.fail-fast applies to the entire matrix. If jobs.
<job_id>.strategy.fail-fast is set to true or its expression evaluates to true, GitHub will cancel all in-progress and queued jobs in the matrix if any job in the matrix fails. This property defaults to true.

You can control how job failures are handled with jobs.<job_id>.strategy.fail-fast and jobs.<job_id>.continue-on-error.

jobs.<job_id>.strategy.fail-fast applies to the entire matrix. If jobs.
<job_id>.strategy.fail-fast is set to true or its expression evaluates to true, GitHub will cancel all in-progress and queued jobs in the matrix if any job in the matrix fails. This property defaults to true.

jobs.<job_id>.continue-on-error applies to a single job. If jobs.<job_id>.continue-on-error is true, other jobs in the matrix will continue running even if the job with jobs.
<job_id>.continue-on-error: true fails.

You can use <code>jobs.<job_id>.strategy.fail-fast</code> and <code>jobs.<job_id>.continue-on-error</code> together. For example, the following workflow will start four jobs. For each job, <code>continue-on-error</code> is determined by the value of <code>matrix.experimental</code>. If any of the jobs with <code>continue-on-error</code>: false fail, all jobs that are in progress or queued will be cancelled. If the job with <code>continue-on-error</code>: true fails, the other jobs will not be affected.

jobs.<job_id>.strategy.max-parallel ₽

By default, GitHub will maximize the number of jobs run in parallel depending on runner availability.

jobs.<job_id>.continue-on-error ∂

jobs.<job_id>.continue-on-error applies to a single job. If jobs.<job_id>.continue-on-error is true, other jobs in the matrix will continue running even if the job with jobs.
<job_id>.continue-on-error: true fails.

Prevents a workflow run from failing when a job fails. Set to true to allow a workflow run to pass when this job fails.

Example: Preventing a specific failing matrix job from failing a workflow run &

You can allow specific jobs in a job matrix to fail without failing the workflow run. For example, if you wanted to only allow an experimental job with node set to 15 to fail without failing the workflow run.

```
runs-on: ${{ matrix.os }}
continue-on-error: ${{ matrix.experimental }}
strategy:
    fail-fast: false
    matrix:
    node: [13, 14]
    os: [macos-latest, ubuntu-latest]
    experimental: [false]
    include:
        - node: 15
        os: ubuntu-latest
        experimental: true
```

jobs.<job_id>.container ∂

Note

If your workflows use Docker container actions, job containers, or service containers, then you must use a Linux runner:

• If you are using GitHub-hosted runners, you must use an Ubuntu runner.

• If you are using self-hosted runners, you must use a Linux machine as your runner and Docker must be installed.

Use jobs.<job_id>.container to create a container to run any steps in a job that don't already specify a container. If you have steps that use both script and container actions, the container actions will run as sibling containers on the same network with the same volume mounts.

If you do not set a container, all steps will run directly on the host specified by runs-on unless a step refers to an action configured to run in a container.

Note

The default shell for run steps inside a container is sh instead of bash. This can be overridden with jobs.<job_id>.defaults.run or jobs.<job_id>.steps[*].shell.

Example: Running a job within a container *∂*

```
YAML
                                                                                    Q
 name: CI
 on:
   push:
     branches: [ main ]
 jobs:
   container-test-job:
     runs-on: ubuntu-latest
     container:
       image: node:18
       env:
        NODE_ENV: development
       ports:
         - 80
       volumes:
         - my_docker_volume:/volume_mount
       options: --cpus 1
     steps:
       - name: Check for dockerenv file
         run: (ls /.dockerenv && echo Found dockerenv) || (echo No dockerenv)
```

When you only specify a container image, you can omit the image keyword.

```
jobs:
  container-test-job:
```

runs-on: ubuntu-latest
container: node:18

jobs.<job_id>.container.image ∂

Use jobs.<job_id>.container.image to define the Docker image to use as the container to run the action. The value can be the Docker Hub image name or a registry name.

Note

Docker Hub normally imposes rate limits on both push and pull operations which will affect jobs on self-hosted runners. However, GitHub-hosted runners are not subject to these limits based on an agreement between GitHub and Docker.

jobs.<job_id>.container.credentials ₽

If the image's container registry requires authentication to pull the image, you can use jobs. <job_id>.container.credentials to set a map of the username and password. The credentials are the same values that you would provide to the docker login command.

Example: Defining credentials for a container registry \mathscr{D}

```
container:
  image: ghcr.io/owner/image
  credentials:
    username: ${{ github.actor }}
    password: ${{ secrets.github_token }}
```

jobs.<job_id>.container.env ∂

Use jobs.<job_id>.container.env to set a map of environment variables in the container.

jobs.<job_id>.container.ports ₽

Use jobs.<job_id>.container.ports to set an array of ports to expose on the container.

jobs.<job_id>.container.volumes ∂

Use jobs.<job_id>.container.volumes to set an array of volumes for the container to use. You can use volumes to share data between services or other steps in a job. You can specify named Docker volumes, anonymous Docker volumes, or bind mounts on the host.

To specify a volume, you specify the source and destination path:

<source>:<destinationPath> .

The <source> is a volume name or an absolute path on the host machine, and <destinationPath> is an absolute path in the container.

Example: Mounting volumes in a container *∂*

volumes:

- my_docker_volume:/volume_mount
- /data/my data
- /source/directory:/destination/directory

jobs.<job_id>.container.options ∂

Use jobs.<job_id>.container.options to configure additional Docker container resource options. For a list of options, see docker create options.

The --network and --entrypoint options are not supported.

jobs.<job_id>.services ∂

(i) Note

If your workflows use Docker container actions, job containers, or service containers, then you must use a Linux runner:

• If you are using GitHub-hosted runners, you must use an Ubuntu runner.

• If you are using self-hosted runners, you must use a Linux machine as your runner and Docker must be installed.

Used to host service containers for a job in a workflow. Service containers are useful for creating databases or cache services like Redis. The runner automatically creates a Docker network and manages the life cycle of the service containers.

If you configure your job to run in a container, or your step uses container actions, you don't need to map ports to access the service or action. Docker automatically exposes all ports between containers on the same Docker user-defined bridge network. You can directly reference the service container by its hostname. The hostname is automatically mapped to the label name you configure for the service in the workflow.

If you configure the job to run directly on the runner machine and your step doesn't use a container action, you must map any required Docker service container ports to the Docker host (the runner machine). You can access the service container using localhost and the mapped port.

For more information about the differences between networking service containers, see Communicating with Docker service containers.

Example: Using localhost *∂*

This example creates two services: nginx and redis. When you specify the container port but not the host port, the container port is randomly assigned to a free port on the host. GitHub sets the assigned host port in the \${{job.services.<service_name>.ports}} context. In this example, you can access the service host ports using the \${{ job.services.nginx.ports['80'] }} and \${{ job.services.redis.ports['6379'] }} contexts.

```
services:
    nginx:
    image: nginx
    # Map port 8080 on the Docker host to port 80 on the nginx container
    ports:
        - 8080:80
    redis:
        image: redis
        # Map random free TCP port on Docker host to port 6379 on redis container
        ports:
        - 6379/tcp
steps:
        - run: |
        echo "Redis available on 127.0.0.1:${{ job.services.redis.ports['6379'] }}"
        echo "Nginx available on 127.0.0.1:${{ job.services.nginx.ports['80'] }}"
```

jobs.<job_id>.services.<service_id>.image ∂

The Docker image to use as the service container to run the action. The value can be the Docker Hub image name or a registry name.

If jobs.<job_id>.services.<service_id>.image is assigned an empty string, the service will not start. You can use this to set up conditional services, similar to the following example.

```
services:
  nginx:
  image: ${{ options.nginx == true && 'nginx' || '' }}
```

jobs.<job_id>.services.<service_id>.credentials



If the image's container registry requires authentication to pull the image, you can use <code>jobs.</code> <job_id>.container.credentials to set a map of the username and password. The credentials are the same values that you would provide to the <code>docker login</code> command.

Example of jobs.<job_id>.services.<service_id>.credentials &

```
services:
    myservice1:
    image: ghcr.io/owner/myservice1
    credentials:
        username: ${{ github.actor }}
        password: ${{ secrets.github_token }}

myservice2:
    image: dockerhub_org/myservice2
    credentials:
        username: ${{ secrets.DOCKER_USER }}
        password: ${{ secrets.DOCKER_PASSWORD }}
```

jobs.<job_id>.services.<service_id>.env @

Sets a map of environment variables in the service container.

jobs.<job_id>.services.<service_id>.ports ♂

Sets an array of ports to expose on the service container.

jobs.<job_id>.services.<service_id>.volumes @

Sets an array of volumes for the service container to use. You can use volumes to share data between services or other steps in a job. You can specify named Docker volumes, anonymous Docker volumes, or bind mounts on the host.

To specify a volume, you specify the source and destination path:

<source>:<destinationPath> .

The <source> is a volume name or an absolute path on the host machine, and <destinationPath> is an absolute path in the container.

Example of jobs.<job_id>.services.<service_id>.volumes ∂

volumes:

- my_docker_volume:/volume_mount
- /data/my_data
- /source/directory:/destination/directory

jobs.<job_id>.services.<service_id>.options ∂

Additional Docker container resource options. For a list of options, see docker create options.

Warning

The --network option is not supported.

jobs.<job_id>.uses ∂

The location and version of a reusable workflow file to run as a job. Use one of the following syntaxes:

- {owner}/{repo}/.github/workflows/{filename}@{ref} for reusable workflows in public and private repositories.
- ./.github/workflows/{filename} for reusable workflows in the same repository.

In the first option, {ref} can be a SHA, a release tag, or a branch name. If a release tag and a branch have the same name, the release tag takes precedence over the branch name. Using the commit SHA is the safest option for stability and security. For more information, see Secure use reference.

If you use the second syntax option (without {owner}/{repo} and @{ref}) the called workflow is from the same commit as the caller workflow. Ref prefixes such as refs/heads and refs/tags are not allowed. You cannot use contexts or expressions in this keyword.

Example of jobs.<job_id>.uses ∂

```
jobs:
    call-workflow-1-in-local-repo:
        uses: octo-org/this-repo/.github/workflows/workflow-
1.yml@172239021f7ba04fe7327647b213799853a9eb89
    call-workflow-2-in-local-repo:
        uses: ./.github/workflows/workflow-2.yml
    call-workflow-in-another-repo:
        uses: octo-org/another-repo/.github/workflows/workflow.yml@v1
```

For more information, see Reuse workflows.

jobs.<job_id>.with ∂

When a job is used to call a reusable workflow, you can use with to provide a map of inputs that are passed to the called workflow.

Any inputs that you pass must match the input specifications defined in the called workflow.

Unlike jobs.<job_id>.steps[*].with, the inputs you pass with jobs.<job_id>.with are not available as environment variables in the called workflow. Instead, you can reference the inputs by using the inputs context.

Example of jobs.<job_id>.with ∂

```
jobs:
   call-workflow:
    uses: octo-org/example-repo/.github/workflows/called-workflow.yml@main
    with:
        username: mona
```

jobs.<job_id>.with.<input_id> @

A pair consisting of a string identifier for the input and the value of the input. The identifier must match the name of an input defined by on.workflow_call.inputs.<inputs_id> in the called workflow. The data type of the value must match the type defined by on.workflow_call.inputs.
<input_id>.type in the called workflow.

Allowed expression contexts: github, and needs.

jobs.<job_id>.secrets ∂

When a job is used to call a reusable workflow, you can use secrets to provide a map of secrets that are passed to the called workflow.

Any secrets that you pass must match the names defined in the called workflow.

Example of jobs.<job id>.secrets ℰ

```
jobs:
   call-workflow:
   uses: octo-org/example-repo/.github/workflows/called-workflow.yml@main
   secrets:
   access-token: ${{ secrets.PERSONAL_ACCESS_TOKEN }}
```

jobs.<job_id>.secrets.inherit ♂

Use the inherit keyword to pass all the calling workflow's secrets to the called workflow. This includes all secrets the calling workflow has access to, namely organization, repository, and environment secrets. The inherit keyword can be used to pass secrets across repositories within the same organization, or across organizations within the same enterprise.

Example of jobs.<job_id>.secrets.inherit @

```
on:
    workflow_dispatch:

jobs:
    pass-secrets-to-workflow:
        uses: ./.github/workflows/called-workflow.yml
        secrets: inherit
```

```
on:
    workflow_call:

jobs:
    pass-secret-to-action:
    runs-on: ubuntu-latest
    steps:
    - name: Use a repo or org secret from the calling workflow.
    run: echo ${{ secrets.CALLING_WORKFLOW_SECRET }}
```

jobs.<job_id>.secrets.<secret_id> ∂

A pair consisting of a string identifier for the secret and the value of the secret. The identifier must match the name of a secret defined by <a href="mailto:on.workflow_call.secrets.<secret_id">on.workflow_call.secrets.<secret_id in the called workflow.

Allowed expression contexts: github, needs, and secrets.

Filter pattern cheat sheet *∂*

You can use special characters in path, branch, and tag filters.

- *: Matches zero or more characters, but does not match the / character. For example,
 0cto* matches 0ctocat.
- **: Matches zero or more of any character.
- ?: Matches zero or one of the preceding character.
- +: Matches one or more of the preceding character.

- [] Matches one alphanumeric character listed in the brackets or included in ranges. Ranges can only include a-z , A-Z , and 0-9 . For example, the range [0-9a-z] matches any digit or lowercase letter. For example, [CB]at matches Cat or Bat and [1-2]00 matches 100 and 200 .
- ! : At the start of a pattern makes it negate previous positive patterns. It has no special meaning if not the first character.

The characters * , [, and ! are special characters in YAML. If you start a pattern with * , [, or ! , you must enclose the pattern in quotes. Also, if you use a <u>flow sequence</u> with a pattern containing [and/or] , the pattern must be enclosed in quotes.

```
# Valid
paths:
    - '**/README.md'

# Invalid - creates a parse error that
# prevents your workflow from running.
paths:
    - ***/README.md

# Valid
branches: [ main, 'release/v[0-9].[0-9]' ]

# Invalid - creates a parse error
branches: [ main, release/v[0-9].[0-9] ]
```

For more information about branch, tag, and path filter syntax, see <a href="https://on.spuths.com/o

Patterns to match branches and tags &

Pattern	Description	Example matches
feature/*	The * wildcard matches any character, but does not match	feature/my-branch
	slash (/).	feature/your-branch
character includi	The ** wildcard matches any character including slash (/) in	feature/beta-a/my-branch
	branch and tag names.	feature/your-branch
		feature/mona/the/octocat

Pattern	Description	Example matches
main	Matches the exact name of a branch or tag name.	main
releases/mona-the-octocat		releases/mona-the-octocat
1*1	Matches all branch and tag names that don't contain a slash	main
	(/). The * character is a special character in YAML. When you start a pattern with * , you must use quotes.	releases
1**1	Matches all branch and tag	all/the/branches
	behavior when you don't use a branches or tags filter.	every/tag
'*feature'	The * character is a special character in YAML. When you	mona-feature
	start a pattern with *, you must use quotes.	feature
		ver-10-feature
v2*	Matches branch and tag names that start with v2.	v2
	that start with V2.	v2.0
		v2.9
v[12].[0-9]+.[0-9]+	Matches all semantic versioning branches and tags with major	v1.10.1
	version 1 or 2.	v2.0.0

Patterns to match file paths \mathscr{D}

Path patterns must match the whole path, and start from the repository's root.

Pattern	Description of matches	Example matches
1*1	The * wildcard matches any character, but does not match	README.md

Pattern	Description of matches	Example matches
	slash (/). The * character is a special character in YAML. When you start a pattern with * , you must use quotes.	server.rb
'*.jsx?'	The ? character matches zero or one of the preceding character.	page.js
		page.jsx
'** [']	The ** wildcard matches any character including slash (/). This is the default behavior when you don't use a path filter.	all/the/files.md
'*.js'	The * wildcard matches any	app.js
	character, but does not match slash (/). Matches all .js files at the root of the repository.	index.js
'**.js'	Matches all .js files in the repository.	index.js
	. 5,5555. ,	js/index.js
		<pre>src/js/app.js</pre>
docs/*	All files within the root of the docs directory only, at the root	docs/README.md
	of the repository.	docs/file.txt
docs/**	Any files in the docs directory and its subdirectories at the root	docs/README.md
	of the repository.	docs/mona/octocat.txt
docs/**/*.md	A file with a .md suffix anywhere	docs/README.md
	in the docs directory.	docs/mona/hello-world.md
		<pre>docs/a/markdown/file.md</pre>
'**/docs/**'	Any files in a docs directory anywhere in the repository.	docs/hello.md
	any interest in the repository.	<pre>dir/docs/my-file.txt</pre>

Pattern	Description of matches	Example matches
		space/docs/plan/space.doc
'**/README.md'	A README.md file anywhere in the repository.	README.md
		js/README.md
'**/*src/**'	Any file in a folder with a src suffix anywhere in the repository.	a/src/app.js
		<pre>my-src/code/js/app.js</pre>
'**/*-post.md'	A file with the suffix -post.md anywhere in the repository.	my-post.md
		path/their-post.md
'**/migrate-*.sql'	A file with the prefix migrate- and suffix .sql anywhere in the	migrate-10909.sql
	repository.	db/migrate-v1.0.sql
		db/sept/migrate-v1.sql
'*.md'	Using an exclamation mark (!) in front of a pattern negates it.	hello.md
'!README.md'	When a file matches a pattern and also matches a negative	Does not match
	pattern defined later in the file, the file will not be included.	README.md
		docs/hello.md
'*.md'	Patterns are checked sequentially. A pattern that	hello.md
'!README.md'	negates a previous pattern will re-include file paths.	README.md
README*		README.doc

Legal

© 2025 GitHub, Inc. <u>Terms Privacy Status Pricing Expert services</u> <u>Blog</u>