

Merge, join, concatenate and compare

pandas provides various methods for combining and comparing `Series` or `DataFrame`.

- `concat()`: Merge multiple `Series` or `DataFrame` objects along a shared index or column
- `DataFrame.join()`: Merge multiple `DataFrame` objects along the columns
- `DataFrame.combine_first()`: Update missing values with non-missing values in the same location
- `merge()`: Combine two `Series` or `DataFrame` objects with SQL-style joining
- `merge_ordered()`: Combine two `Series` or `DataFrame` objects along an ordered axis
- `merge_asof()`: Combine two `Series` or `DataFrame` objects by near instead of exact matching keys
- `Series.compare()` and `DataFrame.compare()`: Show differences in values between two `Series` or `DataFrame` objects

`concat()`

The `concat()` function concatenates an arbitrary amount of `Series` or `DataFrame` objects along an axis while performing optional set logic (union or intersection) of the indexes on the other axes. Like `numpy.concatenate`, `concat()` takes a list or dict of homogeneously-typed objects and concatenates them.

```
In [1]: df1 = pd.DataFrame(  
...:     {  
...:         "A": ["A0", "A1", "A2", "A3"],  
...:         "B": ["B0", "B1", "B2", "B3"],  
...:         "C": ["C0", "C1", "C2", "C3"],  
...:         "D": ["D0", "D1", "D2", "D3"],  
...:     },  
...:     index=[0, 1, 2, 3],  
...: )  
...:
```

[Skip to main content](#)

```

....:     {
....:         "A": ["A4", "A5", "A6", "A7"],
....:         "B": ["B4", "B5", "B6", "B7"],
....:         "C": ["C4", "C5", "C6", "C7"],
....:         "D": ["D4", "D5", "D6", "D7"],
....:     },
....:     index=[4, 5, 6, 7],
....: )
....:

```

```

In [3]: df3 = pd.DataFrame(
....:     {
....:         "A": ["A8", "A9", "A10", "A11"],
....:         "B": ["B8", "B9", "B10", "B11"],
....:         "C": ["C8", "C9", "C10", "C11"],
....:         "D": ["D8", "D9", "D10", "D11"],
....:     },
....:     index=[8, 9, 10, 11],
....: )
....:

```

```
In [4]: frames = [df1, df2, df3]
```

```
In [5]: result = pd.concat(frames)
```

```
In [6]: result
```

```
Out[6]:
```

	A	B	C	D
0	A0	B0	C0	D0
1	A1	B1	C1	D1
2	A2	B2	C2	D2
3	A3	B3	C3	D3
4	A4	B4	C4	D4
5	A5	B5	C5	D5
6	A6	B6	C6	D6
7	A7	B7	C7	D7
8	A8	B8	C8	D8
9	A9	B9	C9	D9
10	A10	B10	C10	D10
11	A11	B11	C11	D11

df1					Result				
	A	B	C	D		A	B	C	D
0	A0	B0	C0	D0	0	A0	B0	C0	D0
1	A1	B1	C1	D1	1	A1	B1	C1	D1
2	A2	B2	C2	D2	2	A2	B2	C2	D2
3	A3	B3	C3	D3	3	A3	B3	C3	D3
df2					4	A4	B4	C4	D4
	A	B	C	D	5	A5	B5	C5	D5
4	A4	B4	C4	D4	6	A6	B6	C6	D6
5	A5	B5	C5	D5	7	A7	B7	C7	D7
6	A6	B6	C6	D6	8	A8	B8	C8	D8
7	A7	B7	C7	D7	9	A9	B9	C9	D9
df3					10	A10	B10	C10	D10
	A	B	C	D	11	A11	B11	C11	D11
8	A8	B8	C8	D8					
9	A9	B9	C9	D9					
10	A10	B10	C10	D10					
11	A11	B11	C11	D11					

Note

`concat()` makes a full copy of the data, and iteratively reusing `concat()` can create unnecessary copies. Collect all `DataFrame` or `Series` objects in a list before using `concat()`.

```
frames = [process_your_file(f) for f in files]
result = pd.concat(frames)
```

Note

When concatenating `DataFrame` with named axes, pandas will attempt to preserve these index/column names whenever possible. In the case where all inputs share a common name, this name will be assigned to the result. When the input names do not all agree, the result will be unnamed. The same is true for `MultiIndex`, but the logic is applied separately on a level-by-level basis.

Joining logic of the resulting axis

[Skip to main content](#)

`join='outer'` takes the union of all axis values

```
In [7]: df4 = pd.DataFrame(
...:     {
...:         "B": ["B2", "B3", "B6", "B7"],
...:         "D": ["D2", "D3", "D6", "D7"],
...:         "F": ["F2", "F3", "F6", "F7"],
...:     },
...:     index=[2, 3, 6, 7],
...: )
...:
```

```
In [8]: result = pd.concat([df1, df4], axis=1)
```

```
In [9]: result
```

```
Out[9]:
```

```

   A  B  C  D  B  D  F
0  A0 B0 C0 D0 NaN NaN NaN
1  A1 B1 C1 D1 NaN NaN NaN
2  A2 B2 C2 D2 B2 D2 F2
3  A3 B3 C3 D3 B3 D3 F3
6  NaN NaN NaN NaN B6 D6 F6
7  NaN NaN NaN NaN B7 D7 F7
```

df1					df4				Result							

`join='inner'` takes the intersection of the axis values

```
In [10]: result = pd.concat([df1, df4], axis=1, join="inner")
```

```
In [11]: result
```

```
Out[11]:
```

```

   A  B  C  D  B  D  F
2  A2 B2 C2 D2 B2 D2 F2
3  A3 B3 C3 D3 B3 D3 F3
```

[Skip to main content](#)

df1					df4				Result							
	A	B	C	D		B	D	F		A	B	C	D	B	D	F
0	A0	B0	C0	D0	2	B2	D2	F2	2	A2	B2	C2	D2	B2	D2	F2
1	A1	B1	C1	D1	3	B3	D3	F3	3	A3	B3	C3	D3	B3	D3	F3
2	A2	B2	C2	D2	6	B6	D6	F6								
3	A3	B3	C3	D3	7	B7	D7	F7								

To perform an effective “left” join using the *exact index* from the original `DataFrame`, result can be reindexed.

```
In [12]: result = pd.concat([df1, df4], axis=1).reindex(df1.index)
```

```
In [13]: result
```

```
Out[13]:
```

	A	B	C	D	B	D	F
0	A0	B0	C0	D0	NaN	NaN	NaN
1	A1	B1	C1	D1	NaN	NaN	NaN
2	A2	B2	C2	D2	B2	D2	F2
3	A3	B3	C3	D3	B3	D3	F3

df1					df4				Result							
	A	B	C	D		B	D	F		A	B	C	D	B	D	F
0	A0	B0	C0	D0	2	B2	D2	F2	0	A0	B0	C0	D0	NaN	NaN	NaN
1	A1	B1	C1	D1	3	B3	D3	F3	1	A1	B1	C1	D1	NaN	NaN	NaN
2	A2	B2	C2	D2	6	B6	D6	F6	2	A2	B2	C2	D2	B2	D2	F2
3	A3	B3	C3	D3	7	B7	D7	F7	3	A3	B3	C3	D3	B3	D3	F3

Ignoring indexes on the concatenation axis

For `DataFrame` objects which don't have a meaningful index, the `ignore_index` ignores overlapping indexes.

```
In [14]: result = pd.concat([df1, df4], ignore_index=True, sort=False)
```

```
In [15]: result
```

```
Out[15]:
```

	A	B	C	D	F
0	A0	B0	C0	D0	NaN
1	A1	B1	C1	D1	NaN
2	A2	B2	C2	D2	NaN
3	A3	B3	C3	D3	NaN

[Skip to main content](#)

6	NaN	B6	NaN	D6	F6
7	NaN	B7	NaN	D7	F7

df1					Result					
	A	B	C	D		A	B	C	D	F
0	A0	B0	C0	D0	0	A0	B0	C0	D0	NaN
1	A1	B1	C1	D1	1	A1	B1	C1	D1	NaN
2	A2	B2	C2	D2	2	A2	B2	C2	D2	NaN
3	A3	B3	C3	D3	3	A3	B3	C3	D3	NaN
df4					4	NaN	B2	NaN	D2	F2
	B	D	F		5	NaN	B3	NaN	D3	F3
2	B2	D2	F2	6	NaN	B6	NaN	D6	F6	
3	B3	D3	F3	7	NaN	B7	NaN	D7	F7	
6	B6	D6	F6							
7	B7	D7	F7							

Concatenating [Series](#) and [DataFrame](#) together

You can concatenate a mix of [Series](#) and [DataFrame](#) objects. The [Series](#) will be transformed to [DataFrame](#) with the column name as the name of the [Series](#).

```
In [16]: s1 = pd.Series(["X0", "X1", "X2", "X3"], name="X")
```

```
In [17]: result = pd.concat([df1, s1], axis=1)
```

```
In [18]: result
```

```
Out[18]:
```

	A	B	C	D	X
0	A0	B0	C0	D0	X0
1	A1	B1	C1	D1	X1
2	A2	B2	C2	D2	X2
3	A3	B3	C3	D3	X3

df1					s1		Result					
	A	B	C	D		X		A	B	C	D	X
0	A0	B0	C0	D0	0	X0	0	A0	B0	C0	D0	X0
1	A1	B1	C1	D1	1	X1	1	A1	B1	C1	D1	X1
2	A2	B2	C2	D2	2	X2	2	A2	B2	C2	D2	X2
3	A3	B3	C3	D3	3	X3	3	A3	B3	C3	D3	X3

Unnamed [Series](#) will be numbered consecutively.

```
In [19]: s2 = pd.Series(["_0", "_1", "_2", "_3"])
```

```
In [20]: result = pd.concat([df1, s2, s2, s2], axis=1)
```

```
In [21]: result
```

```
Out[21]:
```

```

      A  B  C  D  0  1  2
0  A0  B0  C0  D0  _0  _0  _0
1  A1  B1  C1  D1  _1  _1  _1
2  A2  B2  C2  D2  _2  _2  _2
3  A3  B3  C3  D3  _3  _3  _3
```

df1					s2		Result							
	A	B	C	D				A	B	C	D	0	1	2
0	A0	B0	C0	D0	0	_0	0	A0	B0	C0	D0	_0	_0	_0
1	A1	B1	C1	D1	1	_1	1	A1	B1	C1	D1	_1	_1	_1
2	A2	B2	C2	D2	2	_2	2	A2	B2	C2	D2	_2	_2	_2
3	A3	B3	C3	D3	3	_3	3	A3	B3	C3	D3	_3	_3	_3

`ignore_index=True` will drop all name references.

```
In [22]: result = pd.concat([df1, s1], axis=1, ignore_index=True)
```

```
In [23]: result
```

```
Out[23]:
```

```

      0  1  2  3  4
0  A0  B0  C0  D0  X0
1  A1  B1  C1  D1  X1
2  A2  B2  C2  D2  X2
3  A3  B3  C3  D3  X3
```

df1					s1		Result					
	A	B	C	D		X		0	1	2	3	4
0	A0	B0	C0	D0	0	X0	0	A0	B0	C0	D0	X0
1	A1	B1	C1	D1	1	X1	1	A1	B1	C1	D1	X1
2	A2	B2	C2	D2	2	X2	2	A2	B2	C2	D2	X2
3	A3	B3	C3	D3	3	X3	3	A3	B3	C3	D3	X3

Resulting keys

The `keys` argument adds another axis level to the resulting index or column (creating a `MultiIndex`) associate specific keys with each original `DataFrame`.

```
In [24]: result = pd.concat(frames, keys=["x", "y", "z"])
```

```
In [25]: result
```

```
Out[25]:
```

```

      A    B    C    D
x 0  A0  B0  C0  D0
  1  A1  B1  C1  D1
  2  A2  B2  C2  D2
  3  A3  B3  C3  D3
y 4  A4  B4  C4  D4
  5  A5  B5  C5  D5
  6  A6  B6  C6  D6
  7  A7  B7  C7  D7
z 8  A8  B8  C8  D8
  9  A9  B9  C9  D9
 10 A10 B10 C10 D10
 11 A11 B11 C11 D11
```

```
In [26]: result.loc["y"]
```

```
Out[26]:
```

```

      A    B    C    D
4  A4  B4  C4  D4
5  A5  B5  C5  D5
6  A6  B6  C6  D6
7  A7  B7  C7  D7
```


df1					Result					
	A	B	C	D			A	B	C	D
0	A0	B0	C0	D0	x	0	A0	B0	C0	D0
1	A1	B1	C1	D1	x	1	A1	B1	C1	D1
2	A2	B2	C2	D2	x	2	A2	B2	C2	D2
3	A3	B3	C3	D3	x	3	A3	B3	C3	D3
df2					y	4	A4	B4	C4	D4
	A	B	C	D	y	5	A5	B5	C5	D5
4	A4	B4	C4	D4	y	6	A6	B6	C6	D6
5	A5	B5	C5	D5	y	7	A7	B7	C7	D7
6	A6	B6	C6	D6	z	8	A8	B8	C8	D8
7	A7	B7	C7	D7	z	9	A9	B9	C9	D9
df3					z	10	A10	B10	C10	D10
	A	B	C	D	z	11	A11	B11	C11	D11
8	A8	B8	C8	D8						
9	A9	B9	C9	D9						
10	A10	B10	C10	D10						
11	A11	B11	C11	D11						

The `keys` argument can override the column names when creating a new `DataFrame` based on existing `Series`.

```
In [27]: s3 = pd.Series([0, 1, 2, 3], name="foo")
```

```
In [28]: s4 = pd.Series([0, 1, 2, 3])
```

```
In [29]: s5 = pd.Series([0, 1, 4, 5])
```

```
In [30]: pd.concat([s3, s4, s5], axis=1)
```

```
Out[30]:
```

```
   foo  0  1
0    0  0  0
1    1  1  1
2    2  2  4
3    3  3  5
```

```
In [31]: pd.concat([s3, s4, s5], axis=1, keys=["red", "blue", "yellow"])
```

```
Out[31]:
```

```
   red  blue  yellow
0    0    0    0
1    1    1    1
2    2    2    4
3    3    3    5
```

You can also pass a dict to `concat()` in which case the dict keys will be used for the `keys` argument unless other `keys` argument is specified:

[Skip to main content](#)

```
In [32]: pieces = {"x": df1, "y": df2, "z": df3}
```

```
In [33]: result = pd.concat(pieces)
```

```
In [34]: result
```

```
Out[34]:
```

	A	B	C	D
x 0	A0	B0	C0	D0
1	A1	B1	C1	D1
2	A2	B2	C2	D2
3	A3	B3	C3	D3
y 4	A4	B4	C4	D4
5	A5	B5	C5	D5
6	A6	B6	C6	D6
7	A7	B7	C7	D7
z 8	A8	B8	C8	D8
9	A9	B9	C9	D9
10	A10	B10	C10	D10
11	A11	B11	C11	D11

df1					Result					
	A	B	C	D			A	B	C	D
0	A0	B0	C0	D0	x	0	A0	B0	C0	D0
1	A1	B1	C1	D1	x	1	A1	B1	C1	D1
2	A2	B2	C2	D2	x	2	A2	B2	C2	D2
3	A3	B3	C3	D3	x	3	A3	B3	C3	D3
df2					y	4	A4	B4	C4	D4
4	A4	B4	C4	D4	y	5	A5	B5	C5	D5
5	A5	B5	C5	D5	y	6	A6	B6	C6	D6
6	A6	B6	C6	D6	y	7	A7	B7	C7	D7
7	A7	B7	C7	D7	z	8	A8	B8	C8	D8
df3					z	9	A9	B9	C9	D9
8	A8	B8	C8	D8	z	10	A10	B10	C10	D10
9	A9	B9	C9	D9	z	11	A11	B11	C11	D11
10	A10	B10	C10	D10						
11	A11	B11	C11	D11						

```
In [35]: result = pd.concat(pieces, keys=["z", "y"])
```

```
In [36]: result
```

```
Out[36]:
```

	A	B	C	D
--	---	---	---	---

[Skip to main content](#)

	10	A10	B10	C10	D10
	11	A11	B11	C11	D11
y	4	A4	B4	C4	D4
	5	A5	B5	C5	D5
	6	A6	B6	C6	D6
	7	A7	B7	C7	D7

df1					Result					
	A	B	C	D			A	B	C	D
0	A0	B0	C0	D0						
1	A1	B1	C1	D1						
2	A2	B2	C2	D2						
3	A3	B3	C3	D3						
df2					z	8	A8	B8	C8	D8
	A	B	C	D	z	9	A9	B9	C9	D9
4	A4	B4	C4	D4	z	10	A10	B10	C10	D10
5	A5	B5	C5	D5	z	11	A11	B11	C11	D11
6	A6	B6	C6	D6	y	4	A4	B4	C4	D4
7	A7	B7	C7	D7	y	5	A5	B5	C5	D5
df3					y	6	A6	B6	C6	D6
	A	B	C	D	y	7	A7	B7	C7	D7
8	A8	B8	C8	D8						
9	A9	B9	C9	D9						
10	A10	B10	C10	D10						
11	A11	B11	C11	D11						

The `MultiIndex` created has levels that are constructed from the passed keys and the index of the `DataFrame` pieces:

```
In [37]: result.index.levels
Out[37]: FrozenList([['z', 'y'], [4, 5, 6, 7, 8, 9, 10, 11]])
```

`levels` argument allows specifying resulting levels associated with the `keys`

```
In [38]: result = pd.concat(
.....:     pieces, keys=["x", "y", "z"], levels=[["z", "y", "x", "w"]], names=["
.....: )
.....:
```

```
In [39]: result
Out[39]:
```

	A	B	C	D
--	---	---	---	---

[Skip to main content](#)

y	1	A1	B1	C1	D1
	2	A2	B2	C2	D2
	3	A3	B3	C3	D3
	4	A4	B4	C4	D4
z	5	A5	B5	C5	D5
	6	A6	B6	C6	D6
	7	A7	B7	C7	D7
	8	A8	B8	C8	D8
	9	A9	B9	C9	D9
	10	A10	B10	C10	D10
	11	A11	B11	C11	D11

df1					Result					
	A	B	C	D			A	B	C	D
0	A0	B0	C0	D0	x	0	A0	B0	C0	D0
1	A1	B1	C1	D1	x	1	A1	B1	C1	D1
2	A2	B2	C2	D2	x	2	A2	B2	C2	D2
3	A3	B3	C3	D3	x	3	A3	B3	C3	D3
df2					y	4	A4	B4	C4	D4
4	A4	B4	C4	D4	y	5	A5	B5	C5	D5
5	A5	B5	C5	D5	y	6	A6	B6	C6	D6
6	A6	B6	C6	D6	y	7	A7	B7	C7	D7
7	A7	B7	C7	D7	z	8	A8	B8	C8	D8
df3					z	9	A9	B9	C9	D9
8	A8	B8	C8	D8	z	10	A10	B10	C10	D10
9	A9	B9	C9	D9	z	11	A11	B11	C11	D11
10	A10	B10	C10	D10						
11	A11	B11	C11	D11						

```
In [40]: result.index.levels
```

```
Out[40]: FrozenList([[ 'z', 'y', 'x', 'w'], [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]])
```

Appending rows to a `DataFrame`

If you have a `Series` that you want to append as a single row to a `DataFrame`, you can convert the row into a `DataFrame` and use `concat()`

```
In [41]: result = result.append([{"x": "w", "y": 0, "A": "A0", "B": "B0", "C": "C0", "D": "D0"}])
```

[Skip to main content](#)

```
In [42]: result = pd.concat([df1, s2.to_frame().T], ignore_index=True)
```

```
In [43]: result
```

```
Out[43]:
```

```

   A  B  C  D
0  A0 B0 C0 D0
1  A1 B1 C1 D1
2  A2 B2 C2 D2
3  A3 B3 C3 D3
4  X0 X1 X2 X3

```

df1					Result				
	A	B	C	D		A	B	C	D
0	A0	B0	C0	D0	0	A0	B0	C0	D0
1	A1	B1	C1	D1	1	A1	B1	C1	D1
2	A2	B2	C2	D2	2	A2	B2	C2	D2
3	A3	B3	C3	D3	3	A3	B3	C3	D3
s2					4	X0	X1	X2	X3
	A								
	B								
	C								
	D								

merge()

`merge()` performs join operations similar to relational databases like SQL. Users who are familiar with SQL but new to pandas can reference a [comparison with SQL](#).

Merge types

`merge()` implements common SQL style joining operations.

- **one-to-one**: joining two `DataFrame` objects on their indexes which must contain unique values.
- **many-to-one**: joining a unique index to one or more columns in a different `DataFrame`.
- **many-to-many**: joining columns on columns.

[Skip to main content](#)

Note

When joining columns on columns, potentially a many-to-many join, any indexes on the passed `DataFrame` objects **will be discarded**.

For a **many-to-many** join, if a key combination appears more than once in both tables, the `DataFrame` will have the **Cartesian product** of the associated data.

```
In [44]: left = pd.DataFrame(
.....:     {
.....:         "key": ["K0", "K1", "K2", "K3"],
.....:         "A": ["A0", "A1", "A2", "A3"],
.....:         "B": ["B0", "B1", "B2", "B3"],
.....:     }
.....: )
.....:
```

```
In [45]: right = pd.DataFrame(
.....:     {
.....:         "key": ["K0", "K1", "K2", "K3"],
.....:         "C": ["C0", "C1", "C2", "C3"],
.....:         "D": ["D0", "D1", "D2", "D3"],
.....:     }
.....: )
.....:
```

```
In [46]: result = pd.merge(left, right, on="key")
```

```
In [47]: result
```

```
Out[47]:
```

```
   key  A  B  C  D
0  K0  A0 B0 C0 D0
1  K1  A1 B1 C1 D1
2  K2  A2 B2 C2 D2
3  K3  A3 B3 C3 D3
```

left				right				Result					
	key	A	B		key	C	D		key	A	B	C	D
0	K0	A0	B0	0	K0	C0	D0	0	K0	A0	B0	C0	D0
1	K1	A1	B1	1	K1	C1	D1	1	K1	A1	B1	C1	D1
2	K2	A2	B2	2	K2	C2	D2	2	K2	A2	B2	C2	D2
3	K3	A3	B3	3	K3	C3	D3	3	K3	A3	B3	C3	D3

The `how` argument to `merge()` specifies which keys are included in the resulting table. If a key

[Skip to main content](#)

NA. Here is a summary of the **how** options and their SQL equivalent names:

Merge method	SQL Join Name	Description
left	LEFT OUTER JOIN	Use keys from left frame only
right	RIGHT OUTER JOIN	Use keys from right frame only
outer	FULL OUTER JOIN	Use union of keys from both frames
inner	INNER JOIN	Use intersection of keys from both frames
cross	CROSS JOIN	Create the cartesian product of rows of both frames

```
In [48]: left = pd.DataFrame(
.....:     {
.....:         "key1": ["K0", "K0", "K1", "K2"],
.....:         "key2": ["K0", "K1", "K0", "K1"],
.....:         "A": ["A0", "A1", "A2", "A3"],
.....:         "B": ["B0", "B1", "B2", "B3"],
.....:     }
.....: )
.....:

In [49]: right = pd.DataFrame(
.....:     {
.....:         "key1": ["K0", "K1", "K1", "K2"],
.....:         "key2": ["K0", "K0", "K0", "K0"],
.....:         "C": ["C0", "C1", "C2", "C3"],
.....:         "D": ["D0", "D1", "D2", "D3"],
.....:     }
.....: )
.....:

In [50]: result = pd.merge(left, right, how="left", on=["key1", "key2"])

In [51]: result
Out[51]:
```

	key1	key2	A	B	C	D
0	K0	K0	A0	B0	C0	D0
1	K0	K1	A1	B1	NaN	NaN
2	K1	K0	A2	B2	C1	D1
3	K1	K0	A2	B2	C2	D2
4	K2	K1	A3	B3	NaN	NaN

[Skip to main content](#)

left					right					Result						
	key1	key2	A	B		key1	key2	C	D		key1	key2	A	B	C	D
0	K0	K0	A0	B0	0	K0	K0	C0	D0	0	K0	K0	A0	B0	C0	D0
1	K0	K1	A1	B1	1	K1	K0	C1	D1	1	K0	K1	A1	B1	NaN	NaN
2	K1	K0	A2	B2	2	K1	K0	C2	D2	2	K1	K0	A2	B2	C1	D1
3	K2	K1	A3	B3	3	K2	K0	C3	D3	3	K1	K0	A2	B2	C2	D2
										4	K2	K1	A3	B3	NaN	NaN

```
In [52]: result = pd.merge(left, right, how="right", on=["key1", "key2"])
```

```
In [53]: result
```

```
Out[53]:
```

```

key1 key2    A    B    C    D
0    K0    K0    A0    B0    C0    D0
1    K1    K0    A2    B2    C1    D1
2    K1    K0    A2    B2    C2    D2
3    K2    K0    NaN    NaN    C3    D3
```

left					right					Result						
	key1	key2	A	B		key1	key2	C	D		key1	key2	A	B	C	D
0	K0	K0	A0	B0	0	K0	K0	C0	D0	0	K0	K0	A0	B0	C0	D0
1	K0	K1	A1	B1	1	K1	K0	C1	D1	1	K1	K0	A2	B2	C1	D1
2	K1	K0	A2	B2	2	K1	K0	C2	D2	2	K1	K0	A2	B2	C2	D2
3	K2	K1	A3	B3	3	K2	K0	C3	D3	3	K2	K0	NaN	NaN	C3	D3

```
In [54]: result = pd.merge(left, right, how="outer", on=["key1", "key2"])
```

```
In [55]: result
```

```
Out[55]:
```

```

key1 key2    A    B    C    D
0    K0    K0    A0    B0    C0    D0
1    K0    K1    A1    B1    NaN    NaN
2    K1    K0    A2    B2    C1    D1
3    K1    K0    A2    B2    C2    D2
4    K2    K0    NaN    NaN    C3    D3
5    K2    K1    A3    B3    NaN    NaN
```


left					right					Result						
	key1	key2	A	B		key1	key2	C	D		key1	key2	A	B	C	D
0	K0	K0	A0	B0	0	K0	K0	C0	D0	0	K0	K0	A0	B0	C0	D0
1	K0	K1	A1	B1	1	K1	K0	C1	D1	1	K0	K1	A1	B1	NaN	NaN
2	K1	K0	A2	B2	2	K1	K0	C2	D2	2	K1	K0	A2	B2	C1	D1
3	K2	K1	A3	B3	3	K2	K0	C3	D3	3	K1	K0	A2	B2	C2	D2
										4	K2	K0	NaN	NaN	C3	D3
										5	K2	K1	A3	B3	NaN	NaN

```
In [56]: result = pd.merge(left, right, how="inner", on=["key1", "key2"])
```

```
In [57]: result
```

```
Out[57]:
```

```

key1 key2  A  B  C  D
0    K0   K0 A0 B0 C0 D0
1    K1   K0 A2 B2 C1 D1
2    K1   K0 A2 B2 C2 D2

```

left					right					Result						
	key1	key2	A	B		key1	key2	C	D		key1	key2	A	B	C	D
0	K0	K0	A0	B0	0	K0	K0	C0	D0	0	K0	K0	A0	B0	C0	D0
1	K0	K1	A1	B1	1	K1	K0	C1	D1	1	K1	K0	A2	B2	C1	D1
2	K1	K0	A2	B2	2	K1	K0	C2	D2	2	K1	K0	A2	B2	C2	D2
3	K2	K1	A3	B3	3	K2	K0	C3	D3							

```
In [58]: result = pd.merge(left, right, how="cross")
```

```
In [59]: result
```

```
Out[59]:
```

```

key1_x key2_x  A  B key1_y key2_y  C  D
0      K0   K0 A0 B0      K0   K0 C0 D0
1      K0   K0 A0 B0      K1   K0 C1 D1
2      K0   K0 A0 B0      K1   K0 C2 D2
3      K0   K0 A0 B0      K2   K0 C3 D3
4      K0   K1 A1 B1      K0   K0 C0 D0
..      ...   ... ..  ..      ...   ... ..  ..
11     K1   K0 A2 B2      K2   K0 C3 D3
12     K2   K1 A3 B3      K0   K0 C0 D0
13     K2   K1 A3 B3      K1   K0 C1 D1
14     K2   K1 A3 B3      K1   K0 C2 D2
15     K2   K1 A3 B3      K2   K0 C3 D3

```

```
[16 rows x 8 columns]
```

[Skip to main content](#)

left

	key1	key2	A	B
0	K0	K0	A0	B0
1	K0	K1	A1	B1
2	K1	K0	A2	B2
3	K2	K1	A3	B3

right

	key1	key2	C	D
0	K0	K0	C0	D0
1	K1	K0	C1	D1
2	K1	K0	C2	D2
3	K2	K0	C3	D3

Result

	key1_x	key2_x	A	B	key1_y	key2_y	C	D
0	K0	K0	A0	B0	K0	K0	C0	D0
1	K0	K0	A0	B0	K1	K0	C1	D1
2	K0	K0	A0	B0	K1	K0	C2	D2
3	K0	K0	A0	B0	K2	K0	C3	D3
4	K0	K1	A1	B1	K0	K0	C0	D0
5	K0	K1	A1	B1	K1	K0	C1	D1
6	K0	K1	A1	B1	K1	K0	C2	D2
7	K0	K1	A1	B1	K2	K0	C3	D3
8	K1	K0	A2	B2	K0	K0	C0	D0
9	K1	K0	A2	B2	K1	K0	C1	D1
10	K1	K0	A2	B2	K1	K0	C2	D2
11	K1	K0	A2	B2	K2	K0	C3	D3
12	K2	K1	A3	B3	K0	K0	C0	D0
13	K2	K1	A3	B3	K1	K0	C1	D1
14	K2	K1	A3	B3	K1	K0	C2	D2
15	K2	K1	A3	B3	K2	K0	C3	D3

You can [Series](#) and a [DataFrame](#) with a [MultiIndex](#) if the names of the [MultiIndex](#) correspond to the columns from the [DataFrame](#). Transform the [Series](#) to a [DataFrame](#) using [Series.reset_index\(\)](#) before merging

```
In [60]: df = pd.DataFrame({"Let": ["A", "B", "C"], "Num": [1, 2, 3]})

In [61]: df
Out[61]:
   Let  Num
0    A    1
1    B    2
2    C    3

In [62]: ser = pd.Series(
.....:     ["a", "b", "c", "d", "e", "f"],
.....:     index=pd.MultiIndex.from_arrays(
.....:         [ ["A", "B", "C"] * 2, [1, 2, 3, 4, 5, 6]], names=["Let", "Num"]
.....:     ),
.....: )
.....:

In [63]: ser
Out[63]:
```

[Skip to main content](#)

```

B    2    b
C    3    c
A    4    d
B    5    e
C    6    f
dtype: object

```

```
In [64]: pd.merge(df, ser.reset_index(), on=["Let", "Num"])
```

```
Out[64]:
```

```

   Let  Num  0
0    A    1  a
1    B    2  b
2    C    3  c

```

Performing an outer join with duplicate join keys in [DataFrame](#)

```
In [65]: left = pd.DataFrame({"A": [1, 2], "B": [2, 2]})
```

```
In [66]: right = pd.DataFrame({"A": [4, 5, 6], "B": [2, 2, 2]})
```

```
In [67]: result = pd.merge(left, right, on="B", how="outer")
```

```
In [68]: result
```

```
Out[68]:
```

```

   A_x  B  A_y
0     1  2     4
1     1  2     5
2     1  2     6
3     2  2     4
4     2  2     5
5     2  2     6

```

left			right			Result			
	A	B		A	B		A_x	B	A_y
0	1	2	0	4	2	0	1	2	4
1	2	2	1	5	2	1	1	2	5
			2	6	2	2	1	2	6
						3	2	2	4
						4	2	2	5
						5	2	2	6

[Skip to main content](#)

Warning

Merging on duplicate keys significantly increase the dimensions of the result and can cause a memory overflow.

Merge key uniqueness

The `validate` argument checks whether the uniqueness of merge keys. Key uniqueness is checked before merge operations and can protect against memory overflows and unexpected key duplication.

```
In [69]: left = pd.DataFrame({"A": [1, 2], "B": [1, 2]})
In [70]: right = pd.DataFrame({"A": [4, 5, 6], "B": [2, 2, 2]})
In [71]: result = pd.merge(left, right, on="B", how="outer", validate="one_to_one")
-----
MergeError                                Traceback (most recent call last)
Cell In[71], line 1
----> 1 result = pd.merge(left, right, on="B", how="outer", validate="one_to_one")

File ~/work/pandas/pandas/pandas/core/reshape/merge.py:170, in merge(left, right,
    155     return _cross_merge(
    156         left_df,
    157         right_df,
    (...)
    167         copy=copy,
    168     )
    169 else:
--> 170     op = _MergeOperation(
    171         left_df,
    172         right_df,
    173         how=how,
    174         on=on,
    175         left_on=left_on,
    176         right_on=right_on,
    177         left_index=left_index,
    178         right_index=right_index,
    179         sort=sort,
    180         suffixes=suffixes,
    181         indicator=indicator,
    182         validate=validate,
    183     )
    184     return op.get_result(copy=copy)
```

```
File ~/work/pandas/pandas/pandas/core/reshape/merge.py:813, in _MergeOperation.__i
    800 # If argument passed to validate
```

[Skip to main content](#)

```

811 # are in fact unique.
812 if validate is not None:
--> 813     self._validate_validate_kwd(validate)

File ~/work/pandas/pandas/pandas/core/reshape/merge.py:1658, in _MergeOperation._v
1654         raise MergeError(
1655             "Merge keys are not unique in left dataset; not a one-to-one m
1656         )
1657     if not right_unique:
-> 1658         raise MergeError(
1659             "Merge keys are not unique in right dataset; not a one-to-one
1660         )
1662 elif validate in ["one_to_many", "1:m"]:
1663     if not left_unique:

```

MergeError: Merge keys are not unique in right dataset; not a one-to-one merge

If the user is aware of the duplicates in the right `DataFrame` but wants to ensure there are no duplicates in the left `DataFrame`, one can use the `validate='one_to_many'` argument instead, which will not raise an exception.

```

In [72]: pd.merge(left, right, on="B", how="outer", validate="one_to_many")
Out[72]:

```

```

   A_x  B  A_y
0     1  1  NaN
1     2  2  4.0
2     2  2  5.0
3     2  2  6.0

```

Merge result indicator

`merge()` accepts the argument `indicator`. If `True`, a Categorical-type column called `_merge` will be added to the output object that takes on values:

Observation Origin

`_merge` value

Merge key only in `'left'` frame

`left_only`

Merge key only in `'right'` frame

`right_only`

Merge key in both frames

`both`

[Skip to main content](#)

```
In [73]: df1 = pd.DataFrame({"col1": [0, 1], "col_left": ["a", "b"]})
In [74]: df2 = pd.DataFrame({"col1": [1, 2, 2], "col_right": [2, 2, 2]})
In [75]: pd.merge(df1, df2, on="col1", how="outer", indicator=True)
Out[75]:
```

	col1	col_left	col_right	_merge
0	0	a	NaN	left_only
1	1	b	2.0	both
2	2	NaN	2.0	right_only
3	2	NaN	2.0	right_only

A string argument to `indicator` will use the value as the name for the indicator column.

```
In [76]: pd.merge(df1, df2, on="col1", how="outer", indicator="indicator_column")
Out[76]:
```

	col1	col_left	col_right	indicator_column
0	0	a	NaN	left_only
1	1	b	2.0	both
2	2	NaN	2.0	right_only
3	2	NaN	2.0	right_only

Overlapping value columns

The merge `suffixes` argument takes a tuple of list of strings to append to overlapping column names in the input `DataFrame` to disambiguate the result columns:

```
In [77]: left = pd.DataFrame({"k": ["K0", "K1", "K2"], "v": [1, 2, 3]})
In [78]: right = pd.DataFrame({"k": ["K0", "K0", "K3"], "v": [4, 5, 6]})
In [79]: result = pd.merge(left, right, on="k")
In [80]: result
Out[80]:
```

	k	v_x	v_y
0	K0	1	4
1	K0	1	5

left			right			Result			
	k	v		k	v		k	v_x	v_y
0	K0	1	0	K0	4	0	K0	1	4
1	K1	2	1	K0	5	1	K0	1	5
2	K2	3	2	K3	6				

```
In [81]: result = pd.merge(left, right, on="k", suffixes=("_l", "_r"))
```

```
In [82]: result
```

```
Out[82]:
```

```
   k  v_l  v_r
0  K0    1    4
1  K0    1    5
```

left			right			Result			
	k	v		k	v		k	v_l	v_r
0	K0	1	0	K0	4	0	K0	1	4
1	K1	2	1	K0	5	1	K0	1	5
2	K2	3	2	K3	6				

DataFrame.join()

`DataFrame.join()` combines the columns of multiple, potentially differently-indexed `DataFrame` into a single result `DataFrame`.

```
In [83]: left = pd.DataFrame(
.....:     {"A": ["A0", "A1", "A2"], "B": ["B0", "B1", "B2"]}, index=["K0", "K1"]
.....: )
.....:
```

```
In [84]: right = pd.DataFrame(
.....:     {"C": ["C0", "C2", "C3"], "D": ["D0", "D2", "D3"]}, index=["K0", "K2"]
.....: )
.....:
```

```
In [85]: result = left.join(right)
```

```
In [86]: result
```

```
Out[86]:
```

```
   A  B  C  D
K0 A0 B0 C0 D0
K1 A1 B1 NaN NaN
```

[Skip to main content](#)

left			right			Result				
	A	B		C	D		A	B	C	D
K0	A0	B0	K0	C0	D0	K0	A0	B0	C0	D0
K1	A1	B1	K2	C2	D2	K1	A1	B1	NaN	NaN
K2	A2	B2	K3	C3	D3	K2	A2	B2	C2	D2

```
In [87]: result = left.join(right, how="outer")
```

```
In [88]: result
```

```
Out[88]:
```

	A	B	C	D
K0	A0	B0	C0	D0
K1	A1	B1	NaN	NaN
K2	A2	B2	C2	D2
K3	NaN	NaN	C3	D3

left			right			Result				
	A	B		C	D		A	B	C	D
K0	A0	B0	K0	C0	D0	K0	A0	B0	C0	D0
K1	A1	B1	K2	C2	D2	K1	A1	B1	NaN	NaN
K2	A2	B2	K3	C3	D3	K2	A2	B2	C2	D2
						K3	NaN	NaN	C3	D3

```
In [89]: result = left.join(right, how="inner")
```

```
In [90]: result
```

```
Out[90]:
```

	A	B	C	D
K0	A0	B0	C0	D0
K2	A2	B2	C2	D2

left			right			Result				
	A	B		C	D		A	B	C	D
K0	A0	B0	K0	C0	D0	K0	A0	B0	C0	D0
K1	A1	B1	K2	C2	D2	K2	A2	B2	C2	D2
K2	A2	B2	K3	C3	D3					

[Skip to main content](#)

`DataFrame.join()` takes an optional `on` argument which may be a column or multiple column names that the passed `DataFrame` is to be aligned.

```
In [91]: left = pd.DataFrame(
.....:     {
.....:         "A": ["A0", "A1", "A2", "A3"],
.....:         "B": ["B0", "B1", "B2", "B3"],
.....:         "key": ["K0", "K1", "K0", "K1"],
.....:     }
.....: )
.....:
```

```
In [92]: right = pd.DataFrame({"C": ["C0", "C1"], "D": ["D0", "D1"]}, index=["K0",
```

```
In [93]: result = left.join(right, on="key")
```

```
In [94]: result
```

```
Out[94]:
```

```
   A  B key  C  D
0  A0 B0 K0 C0 D0
1  A1 B1 K1 C1 D1
2  A2 B2 K0 C0 D0
3  A3 B3 K1 C1 D1
```

left				right			Result					
	A	B	key		C	D		A	B	key	C	D
0	A0	B0	K0				0	A0	B0	K0	C0	D0
1	A1	B1	K1	K0	C0	D0	1	A1	B1	K1	C1	D1
2	A2	B2	K0	K1	C1	D1	2	A2	B2	K0	C0	D0
3	A3	B3	K1				3	A3	B3	K1	C1	D1

```
In [95]: result = pd.merge(
.....:     left, right, left_on="key", right_index=True, how="left", sort=False
.....: )
.....:
```

```
In [96]: result
```

```
Out[96]:
```

```
   A  B key  C  D
0  A0 B0 K0 C0 D0
1  A1 B1 K1 C1 D1
2  A2 B2 K0 C0 D0
3  A3 B3 K1 C1 D1
```

[Skip to main content](#)

left				right			Result					
	A	B	key		C	D		A	B	key	C	D
0	A0	B0	K0				0	A0	B0	K0	C0	D0
1	A1	B1	K1	K0	C0	D0	1	A1	B1	K1	C1	D1
2	A2	B2	K0	K1	C1	D1	2	A2	B2	K0	C0	D0
3	A3	B3	K1				3	A3	B3	K1	C1	D1

To join on multiple keys, the passed `DataFrame` must have a `MultiIndex`:

```
In [97]: left = pd.DataFrame(
.....:     {
.....:         "A": ["A0", "A1", "A2", "A3"],
.....:         "B": ["B0", "B1", "B2", "B3"],
.....:         "key1": ["K0", "K0", "K1", "K2"],
.....:         "key2": ["K0", "K1", "K0", "K1"],
.....:     }
.....: )

In [98]: index = pd.MultiIndex.from_tuples(
.....:     [("K0", "K0"), ("K1", "K0"), ("K2", "K0"), ("K2", "K1")]
.....: )

In [99]: right = pd.DataFrame(
.....:     {"C": ["C0", "C1", "C2", "C3"], "D": ["D0", "D1", "D2", "D3"]}, index
.....: )

In [100]: result = left.join(right, on=["key1", "key2"])

In [101]: result
Out[101]:
```

	A	B	key1	key2	C	D
0	A0	B0	K0	K0	C0	D0
1	A1	B1	K0	K1	NaN	NaN
2	A2	B2	K1	K0	C1	D1
3	A3	B3	K2	K1	C3	D3

left					right				Result						
	A	B	key1	key2			C	D		A	B	key1	key2	C	D
0	A0	B0	K0	K0	K0	K0	C0	D0	0	A0	B0	K0	K0	C0	D0
1	A1	B1	K0	K1	K1	K0	C1	D1	1	A1	B1	K0	K1	NaN	NaN
2	A2	B2	K1	K0	K2	K0	C2	D2	2	A2	B2	K1	K0	C1	D1
3	A3	B3	K2	K1	K2	K1	C3	D3	3	A3	B3	K2	K1	C3	D3

The default for `DataFrame.join` is to perform a left join which uses only the keys found in the calling `DataFrame`. Other join types can be specified with `how`.

```
In [102]: result = left.join(right, on=["key1", "key2"], how="inner")
```

```
In [103]: result
```

```
Out[103]:
```

```

   A  B key1 key2  C  D
0  A0 B0  K0  K0  C0 D0
2  A2 B2  K1  K0  C1 D1
3  A3 B3  K2  K1  C3 D3
```

left					right				Result						
	A	B	key1	key2			C	D		A	B	key1	key2	C	D
0	A0	B0	K0	K0	K0	K0	C0	D0	0	A0	B0	K0	K0	C0	D0
1	A1	B1	K0	K1	K1	K0	C1	D1	2	A2	B2	K1	K0	C1	D1
2	A2	B2	K1	K0	K2	K0	C2	D2	3	A3	B3	K2	K1	C3	D3
3	A3	B3	K2	K1	K2	K1	C3	D3							

Joining a single Index to a MultiIndex

You can join a `DataFrame` with a `Index` to a `DataFrame` with a `MultiIndex` on a level. The `name` of the `Index` with match the level name of the `MultiIndex`.

```
In [104]: left = pd.DataFrame(
.....:     {"A": ["A0", "A1", "A2"], "B": ["B0", "B1", "B2"]},
.....:     index=pd.Index(["K0", "K1", "K2"], name="key"),
.....: )
```

```
In [105]: index = pd.MultiIndex.from_tuples(
.....:     [("K0", "Y0"), ("K1", "Y1"), ("K2", "Y2"), ("K2", "Y3")],
.....:     names=["key", "Y"],
.....: )
```

[Skip to main content](#)

```
In [106]: right = pd.DataFrame(
.....:     {"C": ["C0", "C1", "C2", "C3"], "D": ["D0", "D1", "D2", "D3"]},
.....:     index=index,
.....: )
.....:
```

```
In [107]: result = left.join(right, how="inner")
```

```
In [108]: result
```

```
Out[108]:
```

```
      A  B  C  D
key Y
K0 Y0 A0 B0 C0 D0
K1 Y1 A1 B1 C1 D1
K2 Y2 A2 B2 C2 D2
     Y3 A2 B2 C3 D3
```

left			right				Result							
		A	B			C	D			A	B	C	D	
K0		A0	B0	K0	Y0	C0	D0	K0	Y0	A0	B0	C0	D0	
K1		A1	B1	K1	Y1	C1	D1	K1	Y1	A1	B1	C1	D1	
K2		A2	B2	K2	Y2	C2	D2	K2	Y2	A2	B2	C2	D2	
				K2	Y3	C3	D3	K2	Y3	A2	B2	C3	D3	

Joining with two MultiIndex

The MultiIndex of the input argument must be completely used in the join and is a subset of the indices in the left argument.

```
In [109]: leftindex = pd.MultiIndex.from_product(
.....:     [list("abc"), list("xy"), [1, 2]], names=["abc", "xy", "num"]
.....: )
.....:
```

```
In [110]: left = pd.DataFrame({"v1": range(12)}, index=leftindex)
```

```
In [111]: left
```

```
Out[111]:
```

```
      v1
abc xy num
a  x   1   0
   2   1
   y   1   2
   2   2
```

[Skip to main content](#)

```

      2      5
y      1      6
      2      7
c  x      1      8
      2      9
      y      1      10
      2      11

```

```

In [112]: rightindex = pd.MultiIndex.from_product(
.....:     [list("abc"), list("xy")], names=["abc", "xy"]
.....: )
.....:

```

```

In [113]: right = pd.DataFrame({"v2": [100 * i for i in range(1, 7)]}, index=right

```

```

In [114]: right

```

```

Out[114]:

```

```

      v2
abc xy
a  x   100
   y   200
b  x   300
   y   400
c  x   500
   y   600

```

```

In [115]: left.join(right, on=["abc", "xy"], how="inner")

```

```

Out[115]:

```

```

      v1  v2
abc xy num
a  x   1    0  100
   2    1  100
   y   1    2  200
   2    3  200
b  x   1    4  300
   2    5  300
   y   1    6  400
   2    7  400
c  x   1    8  500
   2    9  500
   y   1   10  600
   2   11  600

```

```

In [116]: leftindex = pd.MultiIndex.from_tuples(
.....:     [("K0", "X0"), ("K0", "X1"), ("K1", "X2")], names=["key", "X"]
.....: )
.....:

```

```

In [117]: left = pd.DataFrame(
.....:     {"A": ["A0", "A1", "A2"], "B": ["B0", "B1", "B2"]}, index=leftindex
.....: )
.....:

```

[Skip to main content](#)

```

.....: [ ("K0", "Y0"), ("K1", "Y1"), ("K2", "Y2"), ("K2", "Y3")], names=["ke
.....: )
.....:

In [119]: right = pd.DataFrame(
.....:     {"C": ["C0", "C1", "C2", "C3"], "D": ["D0", "D1", "D2", "D3"]}, inde
.....: )
.....:

In [120]: result = pd.merge(
.....:     left.reset_index(), right.reset_index(), on=["key"], how="inner"
.....: ).set_index(["key", "X", "Y"])
.....:

In [121]: result
Out[121]:

```

```

      A  B  C  D
key X  Y
K0  X0 Y0  A0  B0  C0  D0
     X1 Y0  A1  B1  C0  D0
K1  X2 Y1  A2  B2  C1  D1

```

left				right				Result							

Merging on a combination of columns and index levels

Strings passed as the `on`, `left_on`, and `right_on` parameters may refer to either column names or index level names. This enables merging `DataFrame` instances on a combination of index levels and columns without resetting indexes.

```

In [122]: left_index = pd.Index(["K0", "K0", "K1", "K2"], name="key1")

In [123]: left = pd.DataFrame(
.....:     {
.....:         "A": ["A0", "A1", "A2", "A3"],
.....:         "B": ["B0", "B1", "B2", "B3"],
.....:         "key2": ["K0", "K1", "K0", "K1"],
.....:     },

```

[Skip to main content](#)

```

.....:

In [124]: right_index = pd.Index(["K0", "K1", "K2", "K2"], name="key1")

In [125]: right = pd.DataFrame(
.....:     {
.....:         "C": ["C0", "C1", "C2", "C3"],
.....:         "D": ["D0", "D1", "D2", "D3"],
.....:         "key2": ["K0", "K0", "K0", "K1"],
.....:     },
.....:     index=right_index,
.....: )
.....:

In [126]: result = left.merge(right, on=["key1", "key2"])

In [127]: result
Out[127]:
   A  B key2  C  D
key1
K0   A0 B0  K0 C0 D0
K1   A2 B2  K0 C1 D1
K2   A3 B3  K1 C3 D3

```

left				right				Result					
	A	B	key2		C	D	key2		A	B	key2	C	D
K0	A0	B0	K0	K0	C0	D0	K0	K0	A0	B0	K0	C0	D0
K0	A1	B1	K1	K1	C1	D1	K0	K1	A2	B2	K0	C1	D1
K1	A2	B2	K0	K2	C2	D2	K0	K2	A3	B3	K1	C3	D3
K2	A3	B3	K1	K2	C3	D3	K1						

Note

When `DataFrame` are joined on a string that matches an index level in both arguments, the index level is preserved as an index level in the resulting `DataFrame`.

Note

When `DataFrame` are joined using only some of the levels of a `MultiIndex`, the extra levels will be dropped from the resulting join. To preserve those levels, use `DataFrame.reset_index()` on those level names to move those levels to columns prior to the join.

[Skip to main content](#)

Joining multiple `DataFrame`

A list or tuple of `:class:`DataFrame`` can also be passed to `join()` to join them together on their indexes.

```
In [128]: right2 = pd.DataFrame({"v": [7, 8, 9]}, index=["K1", "K1", "K2"])
```

```
In [129]: result = left.join([right, right2])
```

left				right			right2		Result							
	A	B	key2		C	D	key2			A	B	key2_x	C	D	key2_y	v
K0	A0	B0	K0	K0	C0	D0	K0		K0	A0	B0	K0	C0	D0	K0	NaN
K0	A1	B1	K1	K1	C1	D1	K0	K1	7							
K1	A2	B2	K0	K2	C2	D2	K0	K1	8	A2	B2	K0	C1	D1	K0	7.0
K2	A3	B3	K1	K2	C3	D3	K1	K2	9	A2	B2	K0	C1	D1	K0	8.0
										A3	B3	K1	C2	D2	K0	9.0
										A3	B3	K1	C3	D3	K1	9.0

`DataFrame.combine_first()`

`DataFrame.combine_first()` update missing values from one `DataFrame` with the non-missing values in another `DataFrame` in the corresponding location.

```
In [130]: df1 = pd.DataFrame(
.....:     [[np.nan, 3.0, 5.0], [-4.6, np.nan, np.nan], [np.nan, 7.0, np.nan]]
.....: )
.....:
```

```
In [131]: df2 = pd.DataFrame([[-42.6, np.nan, -8.2], [-5.0, 1.6, 4]], index=[1, 2])
```

```
In [132]: result = df1.combine_first(df2)
```

```
In [133]: result
Out[133]:
```

```

   0    1    2
0  NaN  3.0  5.0
1 -4.6  NaN -8.2
2 -5.0  7.0  4.0
```

[Skip to main content](#)

df1				df2				Result			
	0	1	2		0	1	2		0	1	2
0	NaN	3.0	5.0	1	-42.6	NaN	-8.2	0	NaN	3.0	5.0
1	-4.6	NaN	NaN	2	-5.0	1.6	4.0	1	-4.6	NaN	-8.2
2	NaN	7.0	NaN					2	-5.0	7.0	4.0

merge_ordered()

`merge_ordered()` combines order data such as numeric or time series data with optional filling of missing data with `fill_method`.

```
In [134]: left = pd.DataFrame(
.....:     {"k": ["K0", "K1", "K1", "K2"], "lv": [1, 2, 3, 4], "s": ["a", "b",
.....:     ]
.....: }
```

```
In [135]: right = pd.DataFrame({"k": ["K1", "K2", "K4"], "rv": [1, 2, 3]})
```

```
In [136]: pd.merge_ordered(left, right, fill_method="ffill", left_by="s")
```

Out[136]:

```
   k  lv  s  rv
0  K0  1.0 a  NaN
1  K1  1.0 a  1.0
2  K2  1.0 a  2.0
3  K4  1.0 a  3.0
4  K1  2.0 b  1.0
5  K2  2.0 b  2.0
6  K4  2.0 b  3.0
7  K1  3.0 c  1.0
8  K2  3.0 c  2.0
9  K4  3.0 c  3.0
10 K1  NaN d  1.0
11 K2  4.0 d  2.0
12 K4  4.0 d  3.0
```

merge_asof()

`merge_asof()` is similar to an ordered left-join except that matches are on the nearest key rather than equal keys. For each row in the `left DataFrame`, the last row in the `right DataFrame` are selected where the `on` key is less than the left's key. Both `DataFrame` must be

[Skip to main content](#)

Optionally an `merge_asof()` can perform a group-wise merge by matching the `by` key in addition to the nearest match on the `on` key.

```
In [137]: trades = pd.DataFrame(
.....:     {
.....:         "time": pd.to_datetime(
.....:             [
.....:                 "20160525 13:30:00.023",
.....:                 "20160525 13:30:00.038",
.....:                 "20160525 13:30:00.048",
.....:                 "20160525 13:30:00.048",
.....:                 "20160525 13:30:00.048",
.....:             ]
.....:         ),
.....:         "ticker": ["MSFT", "MSFT", "GOOG", "GOOG", "AAPL"],
.....:         "price": [51.95, 51.95, 720.77, 720.92, 98.00],
.....:         "quantity": [75, 155, 100, 100, 100],
.....:     },
.....:     columns=["time", "ticker", "price", "quantity"],
.....: )
.....:

In [138]: quotes = pd.DataFrame(
.....:     {
.....:         "time": pd.to_datetime(
.....:             [
.....:                 "20160525 13:30:00.023",
.....:                 "20160525 13:30:00.023",
.....:                 "20160525 13:30:00.030",
.....:                 "20160525 13:30:00.041",
.....:                 "20160525 13:30:00.048",
.....:                 "20160525 13:30:00.049",
.....:                 "20160525 13:30:00.072",
.....:                 "20160525 13:30:00.075",
.....:             ]
.....:         ),
.....:         "ticker": ["GOOG", "MSFT", "MSFT", "MSFT", "GOOG", "AAPL", "GOOG",
.....:         "bid": [720.50, 51.95, 51.97, 51.99, 720.50, 97.99, 720.50, 52.0
.....:         "ask": [720.93, 51.96, 51.98, 52.00, 720.93, 98.01, 720.88, 52.0
.....:     },
.....:     columns=["time", "ticker", "bid", "ask"],
.....: )
.....:

In [139]: trades
Out[139]:
```

		time	ticker	price	quantity
0	2016-05-25	13:30:00.023	MSFT	51.95	75
1	2016-05-25	13:30:00.038	MSFT	51.95	155
2	2016-05-25	13:30:00.048	GOOG	720.77	100
3	2016-05-25	13:30:00.048	GOOG	720.92	100
4	2016-05-25	13:30:00.048	AAPL	98.00	100

[Skip to main content](#)

Out [140]:

		time	ticker	bid	ask
0	2016-05-25	13:30:00.023	G00G	720.50	720.93
1	2016-05-25	13:30:00.023	MSFT	51.95	51.96
2	2016-05-25	13:30:00.030	MSFT	51.97	51.98
3	2016-05-25	13:30:00.041	MSFT	51.99	52.00
4	2016-05-25	13:30:00.048	G00G	720.50	720.93
5	2016-05-25	13:30:00.049	AAPL	97.99	98.01
6	2016-05-25	13:30:00.072	G00G	720.50	720.88
7	2016-05-25	13:30:00.075	MSFT	52.01	52.03

In [141]: `pd.merge_asof(trades, quotes, on="time", by="ticker")`

Out [141]:

		time	ticker	price	quantity	bid	ask
0	2016-05-25	13:30:00.023	MSFT	51.95	75	51.95	51.96
1	2016-05-25	13:30:00.038	MSFT	51.95	155	51.97	51.98
2	2016-05-25	13:30:00.048	G00G	720.77	100	720.50	720.93
3	2016-05-25	13:30:00.048	G00G	720.92	100	720.50	720.93
4	2016-05-25	13:30:00.048	AAPL	98.00	100	NaN	NaN

`merge_asof()` within `2ms` between the quote time and the trade time.In [142]: `pd.merge_asof(trades, quotes, on="time", by="ticker", tolerance=pd.Timed`

Out [142]:

		time	ticker	price	quantity	bid	ask
0	2016-05-25	13:30:00.023	MSFT	51.95	75	51.95	51.96
1	2016-05-25	13:30:00.038	MSFT	51.95	155	NaN	NaN
2	2016-05-25	13:30:00.048	G00G	720.77	100	720.50	720.93
3	2016-05-25	13:30:00.048	G00G	720.92	100	720.50	720.93
4	2016-05-25	13:30:00.048	AAPL	98.00	100	NaN	NaN

`merge_asof()` within `10ms` between the quote time and the trade time and exclude exact matches on time. Note that though we exclude the exact matches (of the quotes), prior quotes **do** propagate to that point in time.

```
In [143]: pd.merge_asof(
.....:     trades,
.....:     quotes,
.....:     on="time",
.....:     by="ticker",
.....:     tolerance=pd.Timedelta("10ms"),
.....:     allow_exact_matches=False,
.....: )
.....:
```

Out [143]:

		time	ticker	price	quantity	bid	ask
0	2016-05-25	13:30:00.023	MSFT	51.95	75	NaN	NaN
1	2016-05-25	13:30:00.038	MSFT	51.95	155	51.97	51.98

[Skip to main content](#)

3	2016-05-25	13:30:00.048	GOOG	720.92	100	NaN	NaN
4	2016-05-25	13:30:00.048	AAPI	98.00	100	NaN	NaN

compare()

The `Series.compare()` and `DataFrame.compare()` methods allow you to compare two `DataFrame` or `Series`, respectively, and summarize their differences.

```
In [144]: df = pd.DataFrame(
.....:     {
.....:         "col1": ["a", "a", "b", "b", "a"],
.....:         "col2": [1.0, 2.0, 3.0, np.nan, 5.0],
.....:         "col3": [1.0, 2.0, 3.0, 4.0, 5.0],
.....:     },
.....:     columns=["col1", "col2", "col3"],
.....: )
.....:
```

```
In [145]: df
```

```
Out[145]:
```

	col1	col2	col3
0	a	1.0	1.0
1	a	2.0	2.0
2	b	3.0	3.0
3	b	NaN	4.0
4	a	5.0	5.0

```
In [146]: df2 = df.copy()
```

```
In [147]: df2.loc[0, "col1"] = "c"
```

```
In [148]: df2.loc[2, "col3"] = 4.0
```

```
In [149]: df2
```

```
Out[149]:
```

	col1	col2	col3
0	c	1.0	1.0